

Namo

Named dimensional data for Ruby.

Namo is a Ruby library for working with multi-dimensional data using named dimensions. It infers dimensions and coordinates from plain arrays of hashes — the same shape you get from databases, CSV files, JSON, and YAML — so there's no reshaping step.

The design rests on a few stances: every hash key is a dimension and none is privileged as a coordinate or value; formulae attach to a Namu alongside data and re-evaluate on each access, appearing as derived dimensions alongside the data dimensions; operators that combine Namos all take Namos and return Namos — as do the subset-returning Enumerable methods (`select`, `reject`, `sort_by`, `uniq`, and the rest) — so analytical pipelines close; and the formula mechanism is type-agnostic — strings, dates, booleans, and arbitrary Ruby objects work as readily as numbers.

Installation

```
gem install namo
```

Or in your Gemfile:

```
gem 'namo'
```

Usage

Create a Namu instance from an array of hashes:

```
require 'namo'

sales = Namu.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150},
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40},
  {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
])
```

Data may be passed positionally, as above, or by the `data:` keyword where that reads more explicitly:

```
sales = Namu.new(data: [
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100}
])
```

When both are given, the positional argument wins and the keyword `data:` is ignored.

Dimensions and coordinates are inferred:

```
sales.dimensions
# => [:product, :quarter, :price, :quantity]

sales.coordinates[:product]
# => ['Widget', 'Gadget']

sales.coordinates[:quarter]
# => ['Q1', 'Q2']
```

Every key is a dimension; every value is a coordinate. There's no schema declaration and no choosing which column is “the index” — price and quantity are no less first-class than product and quarter.

Selection

Select by named dimension using keyword arguments:

```
# Single value
sales[product: 'Widget']
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150}
# ]>

# Multiple dimensions
sales[product: 'Widget', quarter: 'Q1']
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100}
# ]>

# Range
sales[price: 10.0..20.0]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150}
# ]>

# Array of values
sales[quarter: ['Q1']]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40}
# ]>

# Proc predicate
sales[price: ->(v){v < 20.0}]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150}
# ]>

# Regex predicate
sales[product: /^W/]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150}
# ]>
```

Procs receive the dimension value and select the row when they return truthy. They handle arbitrary predicates — multi-condition tests, nil-aware checks, anything Ruby can express — and compose with everything else:

```
sales[price: ->(v){v < 20.0}, quantity: ->(v){v > 100}]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150}
# ]>
```

Regexes match against the dimension value coerced with `to_s`, so they work against strings, symbols, numbers, dates, or anything else with a sensible string form. `nil` becomes `""` — `//` matches it, `/./` doesn't.

```

sales[product: /widget/i]           # case-insensitive
sales[product: /Widget|Gadget/]     # alternation
sales[product: /^W/, quarter: 'Q1'] # mixed with exact

```

Procs and regexes mix freely with exact values, arrays, ranges, projection, and contraction in the same [] call.

Projection

Project to specific dimensions:

```

sales[:product, :price]
# => #<Namo [
#   {product: 'Widget', price: 10.0},
#   {product: 'Widget', price: 10.0},
#   {product: 'Gadget', price: 25.0},
#   {product: 'Gadget', price: 25.0}
# ]>

```

Selection and projection can be chained:

```

sales[product: 'Widget'][:quarter, :price]
# => #<Namo [
#   {quarter: 'Q1', price: 10.0},
#   {quarter: 'Q2', price: 10.0}
# ]>

```

Or combined in a single call (names before selectors):

```

sales[:quarter, :price, product: 'Widget']
# => #<Namo [
#   {quarter: 'Q1', price: 10.0},
#   {quarter: 'Q2', price: 10.0}
# ]>

```

Contraction

Contraction is the complement of projection. Projection says “keep these dimensions”; contraction says “remove these dimensions, keep everything else”:

```

sales[-:price, -:quantity]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1'},
#   {product: 'Widget', quarter: 'Q2'},
#   {product: 'Gadget', quarter: 'Q1'},
#   {product: 'Gadget', quarter: 'Q2'}
# ]>

```

The -:price syntax uses unary minus on Symbol to produce a negated dimension. Mixing projection and contraction in the same call is an error — the two modes are mutually exclusive:

```

sales[:product, -:price] # => ArgumentError

```

Selection and contraction can be chained:

```

sales[product: 'Widget'][-:price, -:quantity]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1'},
#   {product: 'Widget', quarter: 'Q2'}
# ]>

```

Or combined in a single call (names before selectors):

```
sales[-:price, -:quantity, product: 'Widget']
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1'},
#   {product: 'Widget', quarter: 'Q2'}
# ]>
```

Selection, projection, and contraction always return a new Namo instance, so everything chains.

Concatenation

`+` is the first of Namo's binary operators: it takes a Namo on each side and returns a Namo. The same shape holds for `-`, `&`, `|`, `^`, `==`, `===`, `<`, `<=`, `>`, `>=` and (later) the composition operators — Namo in, Namo (or boolean) out — so analytical pipelines stay queryable end-to-end.

`+` combines two Namo objects that share the same dimensions by appending the rows of the second to the first:

```
q1_sales = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40}
])
```

```
q2_sales = Namo.new([
  {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150},
  {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
])
```

```
all_sales = q1_sales + q2_sales
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40},
#   {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150},
#   {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
# ]>
```

The dimensions must match — the same dimension names in the same order — or the operator raises an `ArgumentError`. Two Namos holding the same columns in a different order must be normalised to a common column order before they can be combined. Formulae carry through from the left-hand side.

Row Removal

`-` removes from the first Namo any row that appears exactly in the second:

```
sales = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150},
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40},
  {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
])
```

```
discontinued = Namo.new([
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40},
  {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
])
```

```

sales - discontinued
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150}
# ]>

```

Removal is exact — every dimension, every value must match. The dimensions must match; different dimensions raise an `ArgumentError`. Formulae carry through from the left-hand side.

Intersection

`&` returns the rows present in both Namo objects, like `Array#&`:

```

sales = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150},
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40},
  {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
])

confirmed = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
])

```

```

sales & confirmed
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Gadget', quarter: 'Q2', price: 25.0, quantity: 60}
# ]>

```

The dimensions must match; different dimensions raise an `ArgumentError`. Formulae carry through from the left-hand side.

Union

`|` returns all rows from both sides, deduplicated, like `Array#|`:

```

q1_sales = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40}
])

all_sales = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Thingy', quarter: 'Q3', price: 5.0, quantity: 10}
])

q1_sales | all_sales
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
#   {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40},
#   {product: 'Thingy', quarter: 'Q3', price: 5.0, quantity: 10}
# ]>

```

The dimensions must match; different dimensions raise an `ArgumentError`. Formulae merge from both sides; the left-hand side's formulae take precedence on conflict.

Symmetric Difference

`^` returns rows that appear in one side but not both:

```
set_a = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40}
])

set_b = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Thingo', quarter: 'Q3', price: 5.0, quantity: 10}
])

set_a ^ set_b
# => #<Namo [
#   {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40},
#   {product: 'Thingo', quarter: 'Q3', price: 5.0, quantity: 10}
# ]>
```

The dimensions must match; different dimensions raise an `ArgumentError`. Formulae merge from both sides; the left-hand side's formulae take precedence on conflict.

Composition

`*` is the equi-join operator. It pairs rows from two `Namos` where coordinates match on every shared dimension, like an inner join on the shared dimension names:

```
ohlcv = Namo.new([
  {symbol: 'BHP', date: '2025-01-01', close: 42.5},
  {symbol: 'RIO', date: '2025-01-01', close: 118.3}
])

fundamentals = Namo.new([
  {symbol: 'BHP', pe: 14.5},
  {symbol: 'RIO', pe: 9.2}
])

ohlcv * fundamentals
# => #<Namo [
#   {symbol: 'BHP', date: '2025-01-01', close: 42.5, pe: 14.5},
#   {symbol: 'RIO', date: '2025-01-01', close: 118.3, pe: 9.2}
# ]>
```

Inner-join semantics: unmatched rows from either side are dropped. Output dimensions are `self.data_dimensions` followed by `other.data_dimensions` exclusive to `other`. Duplicates on shared coordinates are preserved multiplicatively — output multiplicity is the product of input multiplicities on each matching key.

The two `Namos` must have at least one shared data dimension. No overlap raises an `ArgumentError` — the asymmetry with `**` is deliberate, and falling through to a Cartesian product would silently turn a logic error into a large pile of nonsense rows. Formulae merge from both sides; the left-hand side wins on conflict. A name that is data on one side and a formula on the other also raises an `ArgumentError` — the operands disagree about what the name means, with no last-write order to appeal to — so resolve before composing: `audited[-:margin] * modelled`.

Conditional join `*` takes an optional block that decides which matched rows to pair with each left row. Without a block, every shared-dimension match pairs, as above. With one, the block is handed the current

left row and the right rows already matched on the shared dimensions, and returns the subset to pair — the refinement plain `*` can't express, because it pairs every match.

The canonical case is matching each daily price to a single quarterly report — the most recent one dated on or before it. Plain `*` pairs *every* matching quarter; the block narrows that to the one the matching rule picks.

```
prices = Namo.new([
  {symbol: 'BHP', date: '2025-02-15', close: 42.5},
  {symbol: 'BHP', date: '2025-05-20', close: 44.0}
])

quarterly = Namo.new([
  {symbol: 'BHP', quarter_end: '2024-12-31', eps: 1.0},
  {symbol: 'BHP', quarter_end: '2025-03-31', eps: 1.2}
])

prices.*(quarterly) do |row, candidates|
  candidates[quarter_end: ->(qe){qe <= row[:date]}].sort_by{|f| f[:quarter_end]}.last(1)
end
# => #<Namo [
#   {symbol: 'BHP', date: '2025-02-15', close: 42.5, quarter_end: '2024-12-31', eps: 1.0},
#   {symbol: 'BHP', date: '2025-05-20', close: 44.0, quarter_end: '2025-03-31', eps: 1.2}
# ]>
```

`row` is the Row for the current left row, carrying self's formulae, so `row[:date]` and any self formula resolve inside the block. `candidates` is a Namo of the shared-dimension matches, carrying other's formulae, so the block can select on other's derived dimensions too. The block returns a Namo of the rows to pair: one row for the single-match rule above, though it may return zero, one, or many — it's a selector, not a reducer. An empty returned Namo pairs nothing, so that left row is dropped, preserving inner-join semantics. The block can also be passed as a named proc, `prices.*(quarterly, &most_recent_quarter)`.

The block changes only which rows pair. Formulae carry through exactly as in the no-block form — other's merged under self's, self winning on conflict — and the rows the block returns contribute data only.

Cartesian product

`**` is the Cartesian product. Every row from the left paired with every row from the right:

```
products = Namo.new([{:product: 'Widget'}, {:product: 'Gadget'}])
quarters = Namo.new([{:quarter: 'Q1'}, {:quarter: 'Q2'}])

products ** quarters
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1'},
#   {product: 'Widget', quarter: 'Q2'},
#   {product: 'Gadget', quarter: 'Q1'},
#   {product: 'Gadget', quarter: 'Q2'}
# ]>
```

Output has `self.data.length * other.data.length` rows. Output dimensions are `self.data_dimensions + other.data_dimensions`, in operand order. Duplicates are preserved multiplicatively.

The two Namos must have **no** shared data dimensions — the precondition is the mirror image of `*`. Any overlap raises an `ArgumentError`; allowing it would produce rows with the same dimension named twice. Formulae merge from both sides; the left-hand side wins on conflict, and a data/formula name collision between the operands raises, as for `*`.

The visual relationship is intentional: `*` is the filtered version, `**` is the explosive version — more sigil, more output.

Conditional product `**` takes an optional block on the same contract. Where `*`'s block receives the rows pre-matched on the shared dimensions, `**`'s receives *all* of other's rows — there are no shared dimensions to match on — and returns the subset to pair with each left row.

This expresses a conditional product: pair each order with only the shipping tiers that can carry it.

```
orders = Namo.new([
  {order: 'A', weight: 5},
  {order: 'B', weight: 15}
])

tiers = Namo.new([
  {tier: 'light', max_weight: 10},
  {tier: 'heavy', max_weight: 20}
])

orders.**(tiers) do |row, candidates|
  candidates[max_weight: ->(w){w >= row[:weight]}]
end
# => #<Namo [
#   {order: 'A', weight: 5, tier: 'light', max_weight: 10},
#   {order: 'A', weight: 5, tier: 'heavy', max_weight: 20},
#   {order: 'B', weight: 15, tier: 'heavy', max_weight: 20}
# ]>
```

The contract matches `*`'s: `row` carries self's formulae, `candidates` carries other's, the block returns a Namo of rows to pair, and an empty return drops the left row. A block that returns its `candidates` unchanged reproduces the no-block product exactly — `**` is its own block form with the identity selector, just as `*` is `**` with the shared-dimension match applied first.

Decomposition

`/` removes from the left Namo the dimensions that are also in the right, then dedupes the projected rows. It's the inverse of `*` and `**`:

```
combined = Namo.new([
  {symbol: 'BHP', date: '2025-01-01', close: 42.5, pe: 14.5},
  {symbol: 'RIO', date: '2025-01-01', close: 118.3, pe: 9.2}
])

fundamentals = Namo.new([
  {symbol: 'BHP', pe: 14.5},
  {symbol: 'RIO', pe: 9.2}
])

combined / fundamentals
# => #<Namo [
#   {date: '2025-01-01', close: 42.5},
#   {date: '2025-01-01', close: 118.3}
# ]>
```

The intersection of dimensions — here `:symbol` and `:pe` — is removed. Everything else stays. The projected rows are deduplicated, so `/` answers “what's left when these dimensions are factored out?” rather than “what

rows survive a column drop?”. Formulae carry through from the left-hand side.

/ has no precondition. When the two Namos share no dimensions, the intersection is empty, nothing is removed, and self / other returns a Namu equal to self:

```
shipments = Namu.new([order_id: 1, weight: 10])
weather = Namu.new([date: '2025-01-01', temperature: 22])
```

```
shipments / weather
# => #<Namu [order_id: 1, weight: 10]> - equal to shipments
```

The round-trip identity holds for the ** case exactly:

```
a = Namu.new([symbol: 'BHP'], {symbol: 'RIO'})
b = Namu.new([quarter: 'Q1'], {quarter: 'Q2'})
```

```
(a ** b) / b == a
# => true
```

For *, the round-trip is lossy on the dimensions that were shared between the operands:

```
a = Namu.new([symbol: 'BHP', close: 42.5], {symbol: 'RIO', close: 118.3})
b = Namu.new([symbol: 'BHP', pe: 14.5], {symbol: 'RIO', pe: 9.2})
```

```
(a * b) / b
# => #<Namu [close: 42.5], {close: 118.3}>
# Equal to a[-:symbol]. :symbol was shared and is lost.
```

The asymmetry is inherent: / operates only on the two values it receives and can’t distinguish “shared dimension that belonged to both” from “exclusive dimension that belonged only to the right”. Removing the intersection is the only rule expressible from the operands alone, and it gives clean recovery from ** and well-defined (if lossy) recovery from *.

Why / is loose * and ** raise when their preconditions are violated — combining unrelated Namos has no natural answer, and silently producing arbitrary output would turn a logic error into a large pile of nonsense rows. / is different: it’s a projecting operator, not a combining one, and projecting away nothing returns the original. The no-precondition rule isn’t a fallback; it’s the structurally correct result.

This earns / three properties a strict version would lose:

- **Identity test.** combined / other == combined exactly when the two have no shared dimensions — answers “are these Namos dimensionally independent?” without explicit introspection. Same shape as a & b == a answering subset from 0.6.0.
- **Idempotence.** (c / b) / b == c / b. Once b’s dimensions are removed, removing them again does nothing.
- **Pipeline composition.** A processing step that applies / separator can run over any Namu regardless of whether the separator’s dimensions apply. Uninvolved Namos pass through unchanged; involved Namos get stripped. The pipeline doesn’t need to special-case applicability.

This is the same pattern that makes Array#- useful with arrays that aren’t subsets: [1, 2, 3] - [9] == [1, 2, 3], not an error. The no-op-on-non-applicable behaviour lets the operator compose into pipelines that don’t know in advance whether the operation applies.

Equality

Comparison on Namos is **multiset-theoretic on rows**: the comparison operators — ==, eql?, ===, <, <=, >, >= — ignore row order, so two Namos holding the same rows in a different order compare equal, while row multiplicities count (they *are* data). That stance is shared across the equality, pattern-match, and subset/superset operators documented below. Row order is otherwise preserved: to_h and values depend on

it for columnar alignment, and each, first, last, take, drop, and + all observe it. It is the comparison operators alone that treat the sequence of rows as a multiset.

`==` is multiset equality on rows. Class and formulae are ignored; row order is ignored; row multiplicities are not.

```
a = Namo.new([{:x: 1}, {:x: 2}])
b = Namo.new([{:x: 2}, {:x: 1}])
```

```
a == b
# => true
```

```
a == Namo.new([{:x: 1}, {:x: 1}, {:x: 2}])
# => false
```

`eql?` is stricter: it also requires the class to match and the formula names to match. Like `===`, it ignores proc bodies — proc identity isn't a meaningful equivalence in Ruby (`proc{...} == proc{...}` is false), so neither `===` nor `eql?` uses it.

`hash` is consistent with `eql?` and is content-based, so equal Namos hash equally and can be used as Hash keys:

```
h = {a => 'first'}
h[b]
# => 'first'
```

`equal?` is unchanged from Ruby's default — it tests object identity.

`===` answers a different question: does the candidate have the same dimensions and the same formula names? Row data is ignored, and so are the proc bodies themselves — only the names matter. This is the `===` semantics that case statements use, so Namos can serve as templates for analytical shape:

```
sales_shape = Namo.new([{:product: 'X', quarter: 'Q1', price: 0.0, quantity: 0}])
sales_shape[:revenue] = proc{|row| row[:price] * row[:quantity]}
```

```
q1 = Namo.new([{:product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100}])
q1[:revenue] = proc{|row| row[:price] * row[:quantity]}
```

```
sales_shape === q1
# => true (same dimensions, same formula name)
```

```
sales_shape == q1
# => false (different rows)
```

The two `:revenue` procs are independently-written and not the same object — `proc{...} == proc{...}` is false in Ruby. But `===` doesn't compare proc identity; it asks “do these Namos have the same analytical shape?” and the shape is the set of dimensions plus the set of formula names.

Each comparison operator answers a distinct question: `eql?` is strictest (class + data + formula names); `==` is data identity; `===` is analytical identity; the subset operators are data containment.

Rows participate in value semantics on the same data-only basis. `Row#==`, `Row#eql?`, and `Row#hash` compare the underlying row hash and ignore the surrounding Namo's formulae — two Rows with the same data are equal regardless of which Namo yielded them. This makes Rows usable as Hash keys and Set members, and underwrites whole-row deduplication on the Enumerable side:

```
a = Namo.new([{:x: 1}]).first
b = Namo.new([{:x: 1}]).first
a == b           # => true
a.eql?(b)        # => true
```

```
{a => :found}[b] # => :found
```

The omission of `Row#===` and Row-level `</<=>/>=>` is deliberate: a `Row` is a record, not a collection, so the set-theoretic operators don't translate. The value-semantics trio (`==`, `eql?`, `hash`) is what a hash-shaped value needs to behave correctly in Ruby's collection machinery; that's the whole `Row`-comparison story.

Subset and Superset

`<`, `<=`, `>`, `>=` are multiset subset and superset relations on rows.

```
small = Namo.new([{:x: 1}, {:x: 2}])
large = Namo.new([{:x: 1}, {:x: 2}, {:x: 3}])
```

```
small <= large
# => true
```

```
small < large
# => true
```

```
large > small
# => true
```

Equal sets are `<=` and `>=` each other, but neither `<` nor `>`. Disjoint sets are none of the above — unless one side is empty, in which case it is a subset of (and disjoint with) the other.

Multiplicity matters: a single `{x: 1}` is a proper subset of two `{x: 1}`s.

```
one = Namo.new([{:x: 1}])
two = Namo.new([{:x: 1}, {:x: 1}])
```

```
one < two
# => true
```

The dimensions must match; different dimensions raise an `ArgumentError`. Comparing against a non-`Namo` raises a `TypeError`.

Formulae

Define computed dimensions using `[]=:`:

```
sales[:revenue] = proc{|row| row[:price] * row[:quantity]}
```

```
sales[:product, :quarter, :revenue]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', revenue: 1000.0},
#   {product: 'Widget', quarter: 'Q2', revenue: 1500.0},
#   {product: 'Gadget', quarter: 'Q1', revenue: 1000.0},
#   {product: 'Gadget', quarter: 'Q2', revenue: 1500.0}
# ]>
```

Formulae aren't materialised into row data — they re-evaluate on every access. A `:revenue` value reflects the current `:price` and `:quantity` at the moment you ask for it, so derived values stay in sync with whatever the underlying data is doing.

Formulae compose:

```
sales[:cost] = proc{|row| row[:quantity] * 4.0}
sales[:profit] = proc{|row| row[:revenue] - row[:cost]}
```

```

sales[:product, :quarter, :profit]
# => #<Namo [
#   {product: 'Widget', quarter: 'Q1', profit: 600.0},
#   {product: 'Widget', quarter: 'Q2', profit: 900.0},
#   {product: 'Gadget', quarter: 'Q1', profit: 840.0},
#   {product: 'Gadget', quarter: 'Q2', profit: 1260.0}
# ]>

```

Formulae work with selection and projection:

```

sales[product: 'Widget'][:revenue, :quarter]
# => #<Namo [
#   {revenue: 1000.0, quarter: 'Q1'},
#   {revenue: 1500.0, quarter: 'Q2'}
# ]>

```

Formulae carry through selection — a filtered Namo instance remembers its formulae.

Projection of derived dimensions Naming a derived dimension in a projection asks for its values: they are computed against the source and stored in the result's rows, and the formula is dropped — the name is a data dimension of the result. Omitting it carries the formula live, recomputing from the result's own rows on every access:

```

sales[:price, :quantity, :revenue].derived_dimensions
# => [] — :revenue is stored values, a snapshot taken at projection

```

```

sales[:price, :quantity].derived_dimensions
# => [:revenue] — :revenue recomputes from the projected rows on every access

```

The projection list is the selector: name a derived dimension for a snapshot, omit it to keep it as computation. A carried formula whose inputs the projection dropped breaks on access — the same caveat as contracting away a formula's inputs.

Cross-row formulae A formula's arity selects its calling convention. A proc with **one** parameter receives the row, as above. A proc with **two** parameters receives (row, namo), where namo is the Namo the row belongs to — so the formula can reach beyond the current row to the rest of the collection. That's what cross-row computation needs: moving windows, ranks, running totals, anything whose value depends on the row's neighbours.

A simple moving average reads the surrounding rows through namo:

```

prices = Namo.new([
  {symbol: 'AAA', date: 1, close: 10.0},
  {symbol: 'AAA', date: 2, close: 20.0},
  {symbol: 'AAA', date: 3, close: 30.0}
])

prices[:sma] = proc{|row, namo|
  window = namo[symbol: row[:symbol], date: ->(d){d <= row[:date]}]
  window.values(:close).sum / window.count.to_f
}

prices.values(:sma)
# => [10.0, 15.0, 20.0]

```

namo is the Namo that yielded the row, live — so the window always reflects the current state of the object you ask through. A filtered Namo's rows window over the filtered rows; an operator result's rows window

over the result. Appending a row changes every cross-row value on the next access, with no caching.

One-arity formulae are unchanged, and the two forms mix freely — a one-arity formula can reference a two-arity one, and a two-arity formula can reference a one-arity one, by name.

Resolving a two-arity formula needs a `Namo` to window over. A `Row` constructed directly, without one, raises an `ArgumentError` naming the formula rather than letting the missing context surface as an unrelated error.

Parameterised formulae A formula can declare parameters beyond `(row, namo)`. The arguments arrive at access time, through `Row#[]`, so one definition serves every column and every setting:

```
prices[:sma] = proc do |row, namo, field, period|
  window = namo[symbol: row[:symbol], date: ->(d){d <= row[:date]}].last(period)
  window.sum{|r| r[field]} / window.count.to_f
end
```

```
prices.last[:sma, :close, 20]    # 20-period moving average of close
prices.last[:sma, :volume, 50]   # 50-period moving average of volume
```

The number of *required* parameters decides a formula’s calling convention. One means row-scoped, two or more means collection-scoped, and everything past the second receives the arguments given at the call site. A trailing splat or optional after `(row, namo)` makes the arguments optional — `proc{|row, namo, *fields|}` accepts any number, including none. A proc whose second parameter is optional `->(row, namo = nil){...}` requires only one, so it stays row-scoped.

Argument counts are enforced. Asking with the wrong number — too few for the formula’s parameters, too many for a fixed-arity proc, or any at all for a data dimension or an unparameterised formula — raises an `ArgumentError` stating the counts, rather than letting `nil` flow into the formula body:

```
prices.last[:sma]                # ArgumentError: wrong number of arguments for :sma (given 0, expected 2)
prices.last[:close, 20]          # ArgumentError: wrong number of arguments for :close (given 1, expected 0)
```

A formula that requires arguments can’t be materialised without them. `values(:sma)`, `coordinates(:sma)`, naming `:sma` in a projection, and selecting on it all raise the same `ArgumentError`; the no-argument `values`, `coordinates`, and `to_h` omit the dimension, returning everything that can be materialised. `dimensions` and `derived_dimensions` still list it — it is queryable, with arguments. To materialise particular values, bind the arguments in a one-arity wrapper and ask for that:

```
prices[:sma_close_20] = proc{|row| row[:sma, :close, 20]}
prices[:date, :sma_close_20] # materialises per the usual projection rule
```

Polymorphic []=

`[]=` dispatches on the type of the value assigned. A proc registers a formula, as above. Anything else broadcasts the value to every row:

```
sales[:status] = 'active'
sales.values(:status)
# => ['active', 'active', 'active', 'active']
```

```
sales[:revenue] = proc{|row| row[:price] * row[:quantity]}
sales.values(:revenue)
# => [1000.0, 1500.0, 1000.0, 1500.0]
```

The two branches mirror the polymorphism `[]` already has on the selection side, where a single bracket call dispatches over exact values, arrays, ranges, procs, and regexes. Rather than introduce a separate broadcast or `set_all` method for the scalar case, `[]=` reads the same way for both: `sales[:status] = 'active'` says “set status to active across this `Namo`,” and `sales[:revenue] = proc{...}` says “derive revenue from each row.”

The two branches enforce **exclusive storage**: a name is either a data dimension or a derived dimension, never both. Assigning a proc clears any data column of that name; assigning anything else clears any formula of that name. The last write wins, and there is no shadowing:

```
sales[:x] = 5 # :x is a broadcast data value
sales[:x] = proc{|row| row[:price]} # :x is now a formula – the broadcast value is gone
```

```
sales[:x] = proc{|row| row[:price]} # :x is a formula
sales[:x] = 5 # :x is now a broadcast data value – the formula is gone
```

Exclusivity ties directly to the inspection vocabulary: a name assigned a scalar shows up in `data_dimensions`; a name assigned a proc shows up in `derived_dimensions`; never in both, so it appears in dimensions exactly once.

Only a Proc takes the formula branch. An array is a value like any other, so it broadcasts as the per-row value rather than registering as a formula:

```
sales[:weights] = [1, 2, 3]
sales.values(:weights)
# => [[1, 2, 3], [1, 2, 3], [1, 2, 3], [1, 2, 3]]
```

Coordinates and values

`dimensions` covers the *queryable namespace* — every name you can ask for, whether it lives in the row data or is computed by a formula. Once formulae are defined, they appear alongside data dimensions:

```
sales[:revenue] = proc{|row| row[:price] * row[:quantity]}
```

```
sales.dimensions
# => [:product, :quarter, :price, :quantity, :revenue]
```

```
sales.data_dimensions
# => [:product, :quarter, :price, :quantity]
```

```
sales.derived_dimensions
# => [:revenue]
```

`coordinates` gives the unique values per dimension, including derived ones:

```
sales.coordinates[:product]
# => ['Widget', 'Gadget']
```

```
sales.coordinates[:revenue]
# => [1000.0, 1500.0]
```

`values` gives the full per-row sequence — duplicates preserved, row order preserved:

```
sales.values[:product]
# => ['Widget', 'Widget', 'Gadget', 'Gadget']
```

```
sales.values[:revenue]
# => [1000.0, 1500.0, 1000.0, 1500.0]
```

Both `coordinates` and `values` accept positional arguments. With no args they return a Hash across the queryable namespace; with one arg they lazily compute and return just that column as an Array; with multiple args they return a subset Hash containing just the requested columns:

```
sales.values(:product)
# => ['Widget', 'Widget', 'Gadget', 'Gadget']
```

```
sales.values(:product, :quarter)
# => {
#   product: ['Widget', 'Widget', 'Gadget', 'Gadget'],
#   quarter: ['Q1', 'Q2', 'Q1', 'Q2']
# }
```

```
sales.coordinates(:revenue)
# => [1000.0, 1500.0]
```

Single-arg access is lazy: `sales.values(:revenue)` evaluates the formula only across the rows of `:revenue`, without materialising the other columns. The bracket form (`sales.values[:revenue]`) still works through ordinary Hash lookup but pays for the full materialisation up front.

`coordinates` is `values` with `.uniq` applied per column — `coordinates(dim) == values(dim).uniq` holds for every dimension.

`to_h` is the Ruby-conventional alias for the full values Hash:

```
sales.to_h
# => {
#   product: ['Widget', 'Widget', 'Gadget', 'Gadget'],
#   quarter: ['Q1', 'Q2', 'Q1', 'Q2'],
#   price: [10.0, 10.0, 25.0, 25.0],
#   quantity: [100, 150, 40, 60],
#   revenue: [1000.0, 1500.0, 1000.0, 1500.0]
# }
```

Unknown dimensions propagate `nil` per row — `values(:missing)` returns `[nil, nil, ...]` rather than raising or returning a sentinel, matching the convention used by `Row#[]` and `[]` selection. Use `dimensions.include?(:dim)` if you need to check membership directly.

Enumerable

Namo includes `Enumerable`, so `each`, `reduce`, `map`, `select`, `min_by`, and all the rest work out of the box. Rows are yielded as `Row` objects, so formulae are accessible during enumeration:

```
sales.reduce(0){|sum, row| sum + row[:quantity]}
# => 350
```

```
sales[product: 'Widget'].reduce(0){|sum, row| sum + row[:quantity]}
# => 250
```

```
sales[:revenue] = proc{|row| row[:price] * row[:quantity]}
```

```
sales.reduce(0){|sum, row| sum + row[:revenue]}
# => 5000.0
```

```
sales[product: 'Widget'].reduce(0){|sum, row| sum + row[:revenue]}
# => 2500.0
```

```
sales.map{|row| row[:product]}
# => ['Widget', 'Widget', 'Gadget', 'Gadget']
```

```
sales.min_by{|row| row[:price]}[:product]
# => 'Widget'
```

```
sales.flat_map{|row| [row[:price]]}
# => [10.0, 10.0, 25.0, 25.0]
```

The subset-returning Enumerable methods — `select`, `reject`, `sort_by`, `first(n)`, `last(n)`, `take`, `drop`, `take_while`, `drop_while`, `uniq`, and `partition` — return Namos rather than Arrays (`partition` returns `[Namo, Namo]`), carrying formulae through. This keeps the analytical chain closed: the result of a filter is still selectable, projectable, and operable, exactly like the operators that combine Namos:

```
sales.select{|row| row[:price] < 20.0}.values(:price).sum
# => 20.0
```

```
sales.select{|row| row[:price] < 20.0}[product: 'Widget'][:quarter, :revenue].to_a
# => [{quarter: 'Q1', revenue: 1000.0}, {quarter: 'Q2', revenue: 1500.0}]
```

Without an argument, `first` and `last` return a single Row (or `nil` on an empty Namo), following Ruby's convention; with an argument they return a Namo of that many rows. `uniq` dedupes on full-row equality (`Row#==`), or on a block's return value when given one. `select`'s aliases `filter` and `find_all` follow the override and return Namos too.

The transforming and reducing methods are deliberately left as Enumerable's defaults, because their results aren't row-shaped and so can't be a Namo: `map/collect` and `flat_map` return Arrays of whatever the block produces; `reduce/inject`, `sum`, `count`, `min_by`, and `max_by` return scalars. `each` is unchanged — it yields Rows, or returns an Enumerator with no block.

Extracting data

`to_a` returns an array of hashes — the row-oriented form:

```
sales[:product, :quarter, :revenue].to_a
# => [
#   {product: 'Widget', quarter: 'Q1', revenue: 1000.0},
#   {product: 'Widget', quarter: 'Q2', revenue: 1500.0},
#   {product: 'Gadget', quarter: 'Q1', revenue: 1000.0},
#   {product: 'Gadget', quarter: 'Q2', revenue: 1500.0}
# ]
```

`to_h` returns a hash of arrays — the columnar form (see Coordinates and values above):

```
sales[:product, :quarter, :revenue].to_h
# => {
#   product: ['Widget', 'Widget', 'Gadget', 'Gadget'],
#   quarter: ['Q1', 'Q2', 'Q1', 'Q2'],
#   revenue: [1000.0, 1500.0, 1000.0, 1500.0]
# }
```

Named Namos

A Namo can carry a name, passed by the `name:` keyword and read or set through the name accessor:

```
sales = Namo.new(data: rows, name: :sales)
sales.name
# => :sales
```

```
sales.name = :renamed
```

A name defaults to `nil`, and operator results are name-less by design — the result of `+`, `*`, `select`, and the rest is a derived object, not the original, so giving it the parent's name would mislead:

```
(sales + more).name
```



```
# => nil
```

This nil-on-derivation behaviour is what lets subclasses with side effects in `initialize` guard those effects on the name. Operator-derived instances are name-less and skip the side effects; explicitly constructed instances pass `name:` and the side effects fire:

```
class TradingAnalysis < Namo
  def initialize(positional_data = nil, data: [], formulae: {}, name: nil)
    super
    return unless name
    register_indicators
  end
end
```

`super` with no parentheses forwards every argument — positional and keyword — to `Namo#initialize` unchanged. The `return unless name` guard means a subclass need not override every operator to stop the result of `*` or `select from` re-running its construction side effects: it guards on `name` instead.

Why?

Every other multi-dimensional array library requires you to pre-shape your data before you can work with it. `Namo` takes it in the form it likely already comes in.

Name

`Namo`: `nam(ed)` (`dimensi`)`o(ns)`. A companion to `Numo` (numeric arrays for Ruby). And in Aussie culture ‘o’ gets added to the end of names.

Contributing

1. Fork it (<https://github.com/thoran/namo/fork>)
2. Create your feature branch (`git checkout -b my-new-feature`)
3. Commit your changes (`git commit -am 'Add some feature'`)
4. Push to the branch (`git push origin my-new-feature`)
5. Create a new pull request

License

MIT