

Namo Roadmap

Date: 20260615

Design philosophy

Analytical freedom

Namo's foundational design principle is that every key in a row/record/hash is a dimension. There is no distinction between 'axes' and 'values'. This means coordinates are values and values are coordinates.

Namo infers everything from the hash keys. There is no configuration step, no reshaping, no schema declaration. An array of hashes from a database, a CSV, a JSON API, or a YAML file goes straight into Namu and every key is immediately selectable, projectable, and usable in formulae.

This simplification is powerful because it eliminates the distinction that every other tool requires to be made at ingestion time. Analytical decisions belong to the analyst, to be made later during analysis. Every other library in this space requires the user to commit to analytical decisions at ingestion time — which columns are the index, which are data variables, what types they have, when computation runs, where derived values come from. Each of those commitments narrows what the user can later do with the loaded data.

The structure of an analysis — which dimensions are grouped by, which are aggregated over, which serve as identifiers and which as measurements — are determined by the questions asked of the data, not the data as structured. The same dataset answers different questions, and Namu doesn't presage any at ingestion.

A hash {symbol: 'BHP', date: '2025-01-01', close: 42.5} has three dimensions, all equally first-class. You can select on close the same way you select on symbol — `namu[close: 42.5]` is just as valid as `namu[symbol: 'BHP']`. You can project to symbol and date and drop close, or project to close alone. No dimension is privileged. No field is "just data."

`coordinates.close` and `values.close` are the same data viewed two ways — `coordinates` is `values` deduplicated. The shallow API distinction reflects a deep equivalence: every value is a coordinate of its dimension, every coordinate is a value present in the data. Whether you treat `close` as an axis to group by or a quantity to compute over is decided at query time.

Other libraries force the decision earlier. Pandas requires one to choose which columns are the index. xarray requires one to declare dimensions, coordinates, and data variables separately. In each case, the user must understand their data's analytical structure before they can load it — and once loaded, restructuring is awkward.

Type agnosticism

The kinds of computation you'll do are also deferred. Namu doesn't privilege numbers. A formula can concatenate strings, produce status labels, or generate alert messages as naturally as it can compute ratios:

```
e.label = proc{ "#{symbol} ({exchange})" }
e.status = proc{ volume > 0 ? 'active' : 'suspended' }
e.alert = proc{ "#{symbol}: {action} at #{close}" }
```

These are just as valid as `proc{ close / book_value }`. The formula mechanism resolves names and calls procs. What the proc does with the values — arithmetic, string manipulation, date logic, pattern matching — is Ruby's business, not Namu's.

This means Namu can handle datasets that aren't numerical at all: customer records, event logs, text corpora, legal documents with categorised metadata, survey responses. Anything expressible as an array of hashes. The selection, projection, contraction, composition, and set operators all work on text, dates, booleans, and arbitrary objects the same way they work on numbers. `namu[status: 'active']` works because `status` is a dimension like any other, and strings are values like any other.

No other tool in this space offers this. Pandas comes closest — DataFrames can hold mixed types — but its computation model (`df['col'].rolling().mean()`) assumes numeric columns. `xarray` is explicitly numerical. `Quantix` and `Improv` are spreadsheets. `Polars` is a columnar computation engine optimised for numeric and string operations but not arbitrary Ruby objects. `Namo`'s formula mechanism works on whatever Ruby can work on, which is everything.

`Namo` is not a computation engine or a multidimensional spreadsheet. It is a multidimensional database. The distinction matters: spreadsheets and numerical tools like `Quantix`, `Improv`, `xarray`, `Pandas`, and `Polars` are fundamentally organised around numbers. Their dimensions exist to label numerical data. Text in a dimension is a category label, not something you compute on.

Live computation objects

Similarly to how `Namo` treats all dimensions as both data and coordinates, formulae are treated as derived dimensions. `row.close` and `row.earnings_yield` resolve through the same mechanism. The consumer doesn't need to know or care whether a dimension is data or a derived formula. They're all just names with values.

A `Namo` holds data and computation together as a single object whose computed values are always current with respect to its data. Stored values and formulae are accessed through one interface; both reflect current state on every access.

`row.close` returns the stored value. `row.return` invokes the proc and returns the computed value. Same syntax, same semantics from the consumer's perspective. The fact that one is stored and one is computed is an implementation detail of the dimension, not a fact about how it's accessed.

When the underlying data changes — new rows append, existing rows are updated, the backing store delivers different content — every computed value reflects the change without reconciliation. Pull-based reactivity through laziness: each access recomputes from current state, so there's no stale-snapshot problem to solve.

This holds across the formula trajectory: single-row formulae shipped in 0.1.0, cross-row formulae in 0.15.0, parameterised formulae in 0.17.0. Same mechanism throughout, same currentness property regardless of formula complexity.

Other tabular libraries produce snapshots. `Pandas`, `Polars`, `dplyr`, and `DataFrames.jl` all materialise computed values into data columns, severing them from the computation that produced them. Spreadsheets like `Excel` and `Quantix` attach formulae to cells but operate on them through a different interface than data. `Namo` treats them identically across the entire algebra of operations it provides.

Unified treatment of data and derived dimensions

Following from live computation: every operation `Namo` provides treats data dimensions and derived dimensions equivalently. Selection, projection, contraction, composition, set operators, comparison operators, pattern-matching against schemas — none of them distinguish how a dimension's values are produced.

```
namo[:revenue]                # projection – works whether a dimension is data or derived
namo[revenue: 1000..2000]     # selection – works whether a dimension is data or derived
namo - returned_items         # set operation – formulae carry through
namo * fundamentals           # composition – formulae merge by rule
case other; when namo; ...    # pattern-match (Namo#===) – checks data and derived dimensions together
```

Analytical structure can be added or removed without re-ingesting. A user can attach `:return`, `:sma_20`, and `:signal` formulae to an existing `Namo` at any time; subsequent queries treat them as first-class dimensions. They can be removed equally freely. The data layer is untouched; the analytical layer grows and shrinks at the analyst's discretion.

This is what “manipulating derived dimensions as data” means: not a metaphor, but a description of the architecture. Liveness is the mechanism; unified interface is the consequence.

Self-contained formulae

For the unified interface to work, formulae need to be self-contained — pure functions of the row (or (row, namo) for cross-row formulae). They depend on their inputs, not on the environment they were defined in.

```
# Good – pure function of row
e[:return] = proc{|row| row[:close] / row[:previous_close] - 1}

# Bad – depends on external state
@discount_rate = 0.05
e[:discounted] = proc{|row| row[:value] / (1 + @discount_rate)}
```

The self-contained pattern is what makes formulae portable. A Namo can be serialised, sent to another process or another machine, and its formulae will produce the same results because they don't depend on the environment they were defined in. The pattern is also what makes formulae safe to compose — two Namos can be merged by `*` and their formulae coexist because neither references anything outside its inputs.

This isn't enforced by the language (Ruby procs can capture whatever they like), but it's the convention Namo's design depends on. Future versions may move toward stricter enforcement or toward formula representations that don't allow external capture at all.

Algebraic operations on whole Namos

A Namo is a value in an algebra of operations, not just a collection to iterate. Set operators (`+`, `-`, `&`, `|`, `^`) combine Namos as sets of rows. Composition operators (`*`, `**`, `/`) combine Namos with different dimensions. Comparison operators (`==`, `<`, `<=`, `>`, `>=`) ask structural questions about subset relations. `===` asks pattern-match questions about analytical shape.

```
combined = ohlcv * fundamentals    # join on shared dimensions
held      = portfolio & universe    # rows present in both
losers    = today - yesterday       # rows new today
```

These operators take Namos as inputs and produce Namos as outputs. The result of every operator is itself queryable, composable, and operable. The algebra closes.

The shape of an analytical pipeline is not pre-committed to. Operators compose in whatever order makes sense for the question being asked. No other DataFrame library exposes its operations as algebraic operators. Pandas has methods (`merge`, `concat`, `drop_duplicates`); R has function syntax (`bind_rows()`, `inner_join()`); Polars has fluent method chains. Namo treats the operations as first-class operators on the language level, which makes pipelines compact, readable, and uniform.

Composition with surrounding systems

A Namo is passive. It holds state and answers queries. It does not provide subscription mechanisms, event propagation, dependency tracking, or change-detection infrastructure. The reactivity model is the analyst's decision, not Namo's.

Real-time systems built on Namo wire push-based change-detection to the Namo's mutation side and pull-based consumers to the Namo's query side. The Namo doesn't need to know about either; the surrounding system composes them.

```
# Push side: external mechanism updates the Namo
websocket.on_tick{|tick| analysis.append(tick); notifier.broadcast(:updated)}

# Pull side: consumer queries when notified
notifier.subscribe(:updated){|_| update_dashboard(analysis[criteria])}
```

Three components, three responsibilities, clean interfaces. The same passive Namo works in batch reports (pull only, no notifier), interactive analysis (manual queries), polling systems (pull on a timer), and real-time dashboards (push triggers, pull queries) without modification.

Reactive frameworks like RxJS push values through subscription graphs. Namo provides the live computation property that makes such frameworks valuable, but leaves the reactivity infrastructure to the user — to be chosen and composed with whatever the rest of the system already uses. The Namo's job is to be a current-valued data container; everything else is composable around it.

Portability of analytical artefacts

Following from self-contained formulae and the unified interface: a Namo is a portable analytical artefact. Data plus computation plus structure travel together. Serialise a Namo, send it to a colleague, and they have everything needed to query it, extend it, and recompute its derived values.

This is the property notebooks try to provide and fail at. Jupyter notebooks intermix code cells, output cells, and environment state — the artefact you receive is not deterministically re-runnable. A Namo is different: the data is structurally present, the formulae are self-contained procs, and the receiver can call any query operation on it without depending on a particular Ruby environment, package set, or execution order.

This isn't yet a shipped feature — serialisation lands later in 1.x — but the design enables it. Every other principle in this section contributes: self-contained formulae are portable; unified interface means the receiver doesn't need to know what's data vs derived; type agnosticism means the format doesn't have to encode complex type systems; analytical structure at query time means the structure is in the data itself, not in a separate schema.

A Namo is a small, complete, self-describing analytical object. Pandas DataFrame plus the script that produced its computed columns. Excel workbook plus the ability to be queried programmatically. Jupyter notebook minus the bullshit.

Current state: 0.30.0

0.0.0 (2026-03-15): Initial release

Instantiation from an array of hashes. Namo infers dimension names from hash keys and extracts unique coordinate values per dimension automatically. Selection via keyword arguments in [], supporting single values, arrays, and ranges.

```
sales = Namo.new([
  {product: 'Widget', quarter: 'Q1', price: 10.0, quantity: 100},
  {product: 'Widget', quarter: 'Q2', price: 10.0, quantity: 150},
  {product: 'Gadget', quarter: 'Q1', price: 25.0, quantity: 40}
])

sales.dimensions    # => [:product, :quarter, :price, :quantity]
sales.coordinates   # => {product: ['Widget', 'Gadget'], quarter: ['Q1', 'Q2'], ...}
sales[product: 'Widget']      # single value
sales[quarter: ['Q1', 'Q2']]  # array
sales[price: 10.0..20.0]      # range
```

This release established the foundational insight: every key is a dimension, every value is a coordinate. No configuration, no schema, no reshaping step.

0.1.0 (2026-03-28): Formulae and projection

Formulae via []=. A formula is a proc assigned to a name. It receives a Row object and resolves named references against row data or other formulae. Formulae compose — a formula can reference another formula, and the dependency chain resolves lazily through Row.

Projection via positional symbol arguments in []. Selection and projection can be combined in a single call or chained.

```
sales[:revenue] = proc{|r| r[:price] * r[:quantity]}
sales[:profit] = proc{|r| r[:revenue] - r[:cost]}
```

```
sales[:product, :quarter, :revenue]      # projection
sales[:product, :revenue, product: 'Widget'] # projection + selection
sales[product: 'Widget'][:revenue]      # chained
```

Also introduced `to_a` for extracting data as an array of hashes, the `Row` class for formula resolution, and `attr_accessor :data, :formulae`.

0.2.0 (2026-04-14): Enumerable

Namo includes Enumerable. The `each` method yields `Row` objects (not raw hashes), so Enumerable methods like `map`, `select`, `reduce`, `min_by`, and `flat_map` all see formulae as though they were data.

```
sales[:revenue] = proc{|r| r[:price] * r[:quantity]}
total = sales.reduce(0){|sum, row| sum + row[:revenue]}
cheapest = sales.min_by{|row| row[:price]}
```

Selection logic moved from `Namo` to `Row#match?`. Added `Row#to_h`. `each` returns an Enumerator when no block is given.

0.3.0 (2026-04-15): Contraction

Contraction via negated symbols in []. The unary minus on a symbol (`-:price`) produces a `NegatedDimension`, which [] recognises as “remove this dimension from the result.” The complement of projection — projection says what to keep, contraction says what to remove.

```
sales[-:price, -:quantity]      # remove dimensions
sales[-:price, -:quantity, product: 'Widget'] # contraction + selection
```

Raises `ArgumentError` when mixing projection and contraction in the same call. Formulae carry through contraction. `NegatedDimension` and `Symbol#-@` introduced as supporting classes.

`Row` extracted to its own file (`Namo/Row.rb`). Tests split into per-class files.

0.4.0 (2026-04-15): Concatenation and row removal

The first two row-axis set operators.

`+` concatenates two `Namos` with matching dimensions. Appends rows from the second to the first. Formulae merge — self’s formulae take priority on conflict.

`-` removes exact matching rows. Whole-row set difference. If a row appears in the right operand, it’s removed from the left. Both operators raise `ArgumentError` when dimensions differ.

```
jan_sales + feb_sales      # more rows, same dimensions
all_sales - returned_items # fewer rows, same dimensions
```

0.5.0 (2026-04-16): Intersection, union, symmetric difference

The remaining row-axis set operators, completing the set algebra.

`&` returns rows present in both operands. `|` returns all rows deduplicated (implemented via `Array#|`). `^` returns rows in one operand but not both — the symmetric difference.

```

sales & promotions      # rows in both
sales | new_arrivals    # all rows, no duplicates
sales ^ competitor      # rows unique to each side

```

All three require matching dimensions, raise `ArgumentError` otherwise. Formulae carry through from self, merge from other, self wins on conflict. `|` and `^` merge formulae from both sides.

0.6.0 (2026-05-11): Equality, pattern-match, subset, and superset

Equality, pattern-match, and subset/superset operators, multiset-theoretic on rows where row-content matters. Two Namos with the same rows in different orders are equal; two Namos where one's rows are contained in the other are in a subset relation regardless of how either was ordered. Duplicate rows count — `[{a: 1}, {a: 1}]` is not equal to `[{a: 1}]`. This follows from Namo's identity as a multidimensional database rather than a sequence: the comparison operators treat the rows as a multiset, so their order does not affect equality, while row count does. Row order is preserved by the row-sequence operations (`to_h/values`, `each`, `first/last/take/drop`, `+`); it is the comparison operators specifically that disregard it.

The equality hierarchy mirrors Ruby's standard convention, extended with `===` for pattern-match dispatch:

```

a.equal?(b)    # same object (inherited from BasicObject, not overridden)
a.eql?(b)      # same class + multiset-equal data + same formula names
a == b         # multiset-equal data, any class, formula names ignored
a === b        # same dimensions + same formula names, any class, data ignored

```

`eql?` requires class match because subclassing carries included modules (`TradingAnalysis < Namo` includes `Indicators`, `Scoring`); two instances of the same subclass with the same data and formulae are interchangeable as analyses, while a subclass and a bare Namo with the same data are not. `==` ignores class and formulae — two Namos are equal if they hold the same data, mirroring Ruby's `1 == 1.0 / 1.eql?(1.0)` convention.

`===` answers a structurally different question: does this candidate fit this analytical pattern? Same dimensions, same formula names, regardless of the rows. This is what case statements need — pattern dispatch on analytical shape, not on data identity:

```
template = Namo.new([{:x: 0}])
```

```

case Namo.new([{:x: 5}, {:x: 6}])
when template then :matched
else :not_matched
end
# => :matched

```

`===` returns false (rather than raising) for non-Namo operands, since case statements require it to be safe to call on any value. Without a Namo-specific `===`, Ruby's default would delegate to `==`, which would dispatch case statements on multiset row equality — structurally misleading when users want analytical-type dispatch.

```

namo_a = Namo.new([{:symbol: 'BHP', close: 42.5}, {:symbol: 'RIO', close: 118.3}])
namo_b = Namo.new([{:symbol: 'RIO', close: 118.3}, {:symbol: 'BHP', close: 42.5}])

```

```

namo_a == namo_b      # true - same rows, different order
namo_a.eql?(namo_b)   # true - same class, same rows, no formulae either side
namo_a.equal?(namo_b) # false - different objects

```

`hash` is content-based and consistent with `eql?` — computed from the canonical form of `@data` (rows sorted lexicographically), the class, and the formula names. Two Namos that are `eql?` produce the same hash, making Namos usable as Hash keys and Set members.

Subset and superset follow `stdlib` Set's precedent, generalised to multisets:

```
a < b    # a's rows are a strict multiset subset of b's rows
```

```

a <= b    # a's rows are a multiset subset of b's rows
a > b     # a's rows are a strict multiset superset of b's rows
a >= b    # a's rows are a multiset superset of b's rows

```

These pair with the set operators algebraically: $a \& b == a$ iff $a \leq b$; $a | b == b$ iff $a \leq b$; $a - b == \emptyset$ iff $a \leq b$. Dimension mismatch raises `ArgumentError`; non-Namo operand raises `TypeError`.

The error message format for dimension mismatch was standardised to "dimensions don't match: X vs Y" across all binary operators (+, -, &, |, ^, <, <=, >, >=), retrofitted into the set operators from 0.4.0–0.5.0.

Deliberately not included: `<=>` and `Comparable` (subset/superset is a partial order, not total — two Namos can be incomparable, and `stdlib Set` omits `<=>` for the same reason); Row-level public operators; block forms (relaxed-matching variants land in 0.14.0 alongside block forms for set and composition operators).

0.7.0 (2026-05-20): Derived-dimension surfacing, values, live views

Three things land together: values as the third inspection aspect alongside dimensions and coordinates; derived dimensions (formula names) surfacing through the inspection aspects so data and derived are queried uniformly; and explicit accessors for the data/derived split (`data_dimensions`, `derived_dimensions`). All three accessors are live — they recompute from current state on every call, with no memoisation.

The shape of the API is plain Ruby. `dimensions` returns an `Array`, `coordinates` and `values` return `Hash` (eager) or `Array` (single column, lazy), depending on how they're called. No aspect classes, no subclasses of `Array/Hash`. Subclassing of `Namo` itself (`class Sales < Namu`) and `Namo#===` from 0.6.0 already cover case-statement dispatch on analytical shape, so the aspect-class layer is not needed.

values Extract per-dimension sequences, preserving duplicates and row order.

```

namo = Namu.new([
  {symbol: 'BHP', close: 42.5},
  {symbol: 'RIO', close: 118.3},
  {symbol: 'BHP', close: 43.1},
])

namo.values(:symbol)      # => ['BHP', 'RIO', 'BHP']
namo.coordinates(:symbol) # => ['BHP', 'RIO']

```

`coordinates` answers “what exists along this axis?” — the unique axis labels. `values` answers “what does this column actually contain?” — the raw data, preserving count and order. The difference is analogous to `DISTINCT` vs a bare `SELECT` on a column.

Both accept positional dimension arguments:

```

namo.values                # => full Hash {dim => sequence}
namo.values(:symbol)       # => Array – one column, lazily computed
namo.values(:symbol, :close) # => subset Hash with just those columns
namo.values(:unknown)      # => [nil, nil, ...] – one nil per row
namo.values(:symbol, :unknown) # => {symbol: [...], unknown: [nil, nil, ...]}

```

With no args, `values` returns a `Hash` keyed by every dimension in the queryable namespace. With one arg it returns just that column as an `Array`, computing only that column (lazy). With multiple args it returns a `Hash` containing just the requested columns, again computing only those. Unknown dimensions propagate `nil` naturally — `row[:unknown]` is `nil` for each row, so the column is an `Array` of `nil`s — matching the convention used by `Row#[]` and `Namu#[]` selection. `coordinates` mirrors the same shape; an unknown dimension becomes `[nil]` after `uniq`ing the per-row `nil`s.

Asymmetric returns (`Array` for one arg, `Hash` for none or many) follow the natural reading: `values(:symbol)` is “the values of `:symbol`,” singular, so an `Array`. Ruby has precedent for arg-shape-driven returns (`arr[2]` vs `arr[2..3]`); the principle is that the most useful shape for each calling pattern wins over uniformity.

The implementation is straightforward and lives directly in `Namo`:

```
def values(*dims)
  if dims.empty?
    dimensions.each_with_object({}){|dim, hash| hash[dim] = values_for(dim)}
  elsif dims.length == 1
    dim = dims.first
    return nil unless dimensions.include?(dim)
    values_for(dim)
  else
    dims.each_with_object({}){|dim, hash| hash[dim] = values_for(dim) if dimensions.include?(dim)}
  end
end

def coordinates(*dims)
  if dims.empty?
    values.transform_values(&:uniq)
  elsif dims.length == 1
    values(dims.first)&.uniq
  else
    dims.each_with_object({}){|dim, hash| hash[dim] = values(dim)&.uniq if dimensions.include?(dim)}
  end
end
```

`coordinates` is literally `values` with `.uniq` applied per column — a single source of truth for per-dimension column extraction. `coordinates(dim) == values(dim).uniq` holds across the queryable namespace, by construction.

Queryable namespace: data and derived dimensions together `dimensions`, `values`, and `coordinates` all cover the *queryable namespace* — the union of data dimensions (keys of `@data.first`) and derived dimensions (keys of `@formulae`). From the user’s perspective, data and derived are equivalent: anywhere you can name a data dimension, you can name a formula.

```
namo = Namu.new([{:price: 10.0, :quantity: 100}, {:price: 25.0, :quantity: 60}])
namo[:revenue] = proc{|r| r[:price] * r[:quantity]}
```

```
namo.dimensions      # => [:price, :quantity, :revenue]
namo.values(:revenue) # => [1000.0, 1500.0]
namo.coordinates(:revenue) # => [1000.0, 1500.0]
namo.to_h            # => {price: [...], quantity: [...], revenue: [...]}
```

Asking for `values(:revenue)` evaluates the formula across every row and returns the resulting column. The derived dimension behaves identically to a data dimension from the inspection vocabulary’s point of view. This unification matches the 0.1.0 design principle that formulae are first-class dimensions, not a separate computed-column concept; 0.7.0 finishes wiring that principle through the inspection methods.

Data and derived as explicit accessors For the cases where the distinction does matter — set operators need to compare row schemas, projection through `[]` operates on data columns, and code introspecting a `Namo` may want to know which names are data vs derived — the split is exposed as parallel accessors:

```
namo.data_dimensions # => [:price, :quantity]  – keys of the first row
namo.derived_dimensions # => [:revenue]        – keys of @formulae
namo.dimensions      # => [:price, :quantity, :revenue] – the union, queryable namespace
```

The set operators (`+`, `-`, `&`, `|`, `^`) use `data_dimensions` for their matching check internally, preserving the 0.6.0 semantics that two `Namos` with the same data layout can be combined regardless of which formulae

either has registered.

to_h to_h is the Ruby-conventional alias for values with no arguments:

```
def to_h
  values
end
```

Returns the full columnar Hash. values is the primitive name; to_h reads naturally at coercion sites (hash = namo.to_h; hash.keys).

This is purely about output. Internal storage stays row-oriented (@data as Array<Hash>); columnar storage as a performance option is a 3.x concern. The public to_h/values shape is stable from 0.7.0.

The pairing of values with proc-based selection (0.8.0) is about user workflow: if you're writing namo[pe: ->(v){ v && v < 15 }], you probably also want namo.values(:pe) first to inspect the range and spot nils before writing the predicate.

Live views, no memoisation Every call to dimensions, data_dimensions, derived_dimensions, values, coordinates, or to_h recomputes from the Namos current state. Nothing is cached across accesses (within a single values access, 0.29.0 materialises each column once — a loop-invariant hoist, not a cache; see 0.29.0). Mutating @data or @formulae (e.g. via namo[:new_formula] = proc{...}) is reflected immediately on the next inspection call.

```
namo.derived_dimensions # => []
namo[:revenue] = proc{|r| r[:price] * r[:quantity]}
namo.derived_dimensions # => [:revenue]
namo.values(:revenue)   # => [1000.0, 1500.0]
```

The performance cost compared to a memoised implementation is real but unmeasured at typical sizes (hundreds to low thousands of rows). Single-column laziness — values(:dim) computing only the requested column — covers the case where it matters most: a million-row Namos with expensive derived dimensions, where materialising the full Hash would be wasteful. Wholesale aspect caching (freeze-aware materialisation) lands in 2.x as an opt-in, transparent optimisation — see 2.x's caching subsection. It is deliberately not in 1.x: the whole 1.x line is pure-live, recomputing derived views from current state on every access, so that liveness is unconditional and there is no staleness to reason about. The precise invariant is *no derived value outlives the access that produced it* — which permits 0.29.0's within-access materialisation (a column computed once while resolving one access, discarded when it returns) without admitting cross-access caching. Cross-access caching is a 2.x performance concern gated on freeze, never a behaviour change.

Composition through plain types Because dimensions, data_dimensions, and derived_dimensions return plain Array, and values(:dim) and coordinates(:dim) return plain Array, the full vocabulary of Ruby's Array operators applies without Namos defining anything new:

```
a.dimensions == b.dimensions           # exact match on queryable namespace
a.data_dimensions == b.data_dimensions # exact match on data layout
a.derived_dimensions == b.derived_dimensions # exact match on formula names
a.coordinates(:symbol) == b.coordinates(:symbol) # set equality on a dimension
a.values(:close) == b.values(:close) # sequence equality on a dimension
a.values(:close).sum # any Array method
namos.sort_by{|n| n.values(:date).first} # sort Namos by aspect
[:price, :quantity].all?{|d| a.dimensions.include?(d)} # subset-style check
```

This is the answer to “where does ordering live?” — at the aspect level, where projections to plain Arrays inherit <=> and == from Ruby's built-ins. No Namos#<=>, no Namos-level total order. Just plain return types.

Case-statement dispatch: subclassing and `Namo#===` Case-statement dispatch on analytical shape has two paths, both already present without 0.7.0 adding any new `===` infrastructure:

```
# Subclassing – the cleanest path when you have named shapes
class Sales < Namu; end
class Inventory < Namu; end

case incoming
when Sales      then publish_sales_dashboard(incoming)
when Inventory then publish_inventory_dashboard(incoming)
end

# Namu#=== – for ad-hoc shape templates without a subclass
ohlc = Namu.new([{:symbol: nil, date: nil, open: nil, high: nil, low: nil, close: nil, volume: nil}])
case incoming
when ohlc then process_ohlc(incoming)
end
```

`Class#===` dispatches on `is_a?` (stock Ruby), so subclasses of `Namu` flow through case statements directly, inheriting case dispatch alongside their domain methods. `Namu#===` (from 0.6.0) matches on full analytical shape — same dimensions, same formula names — for templates that don't justify a subclass.

Aspect-level `===` (the originally-planned `ohlc.dimensions === incoming` template-match) is not part of 0.7.0. The case-dispatch use it would have served is covered by subclassing for known shapes and `Namu#===` for ad-hoc ones; the data-only and derived-only fine-grained variants don't have a concrete use case beyond what `==` on the relevant accessor already gives (`a.data_dimensions == b.data_dimensions`, etc.). If a case-dispatch need on the finer split materialises later, a small `Matcher` returned by a factory method on `Namu` can serve it without bringing back the aspect-class hierarchy.

Why Array storage, not Set `values` reinforces the decision to keep `@data` as `Array<Hash>` rather than `Set<Hash>`. `values` needs ordering and duplicates — both things `Set` throws away. If the internal store were a `Set`, `values` would lose row order and couldn't contain duplicate rows. `to_a` would need to reconstruct an order that no longer exists internally. `to_h` (columnar) would have the same problem — the column arrays need a consistent row ordering so that `values(:symbol)[i]` and `values(:close)[i]` correspond to the same row. `Set` buys fast membership testing and automatic deduplication, but `Namu` doesn't need either of those as primitives. The set operations (`&`, `|`, `^`) work on Array storage just fine.

0.8.0 (2026-05-21): Proc-based and regex-based selection

Two ways to extend `[]` selection beyond exact values, arrays, and ranges: procs for arbitrary predicates, regexes for string pattern matching. Both are single-branch additions to existing selection logic, paired here because they share the same dispatch site (`Row#match?`).

Proc-based selection `[]` accepts procs as selection predicates on any dimension. The proc receives the dimension value (or `nil` for a missing or nil-valued dimension) and decides — truthy result selects the row.

```
namo[pe: ->(v){v && v < 15}]
namo[price: ->(v){v > 10.0}, symbol: ->(v){v != 'TEST'}]
```

Implementation is a single `when Proc` branch in `Row#match?` calling `coordinate.call(self[dimension])`. The match is on `Proc` specifically — not duck-typed via `respond_to?(:call)` — so lambdas, procs, and `Symbol#to_proc` results all flow through, but methods and other callables don't. Exceptions from the predicate propagate; no rescue clause.

This enables multi-factor screening in one expression:

```
namo[pe: ->(v){v && v < 15}, price_to_book: ->(v){v && v < 1.5}]
```

Proc-based selection composes with contraction and projection in a single `[]` call, and works on formula-defined dimensions:

```
namo[:revenue] = proc{|r| r[:price] * r[:quantity]}
namo[revenue: ->(v){v >= 1500.0}]
```

Regex-based selection `[]` accepts regexes as selection predicates on any dimension. The dimension value is coerced with `to_s` before matching, so regexes work against strings, symbols, integers, floats, dates, and anything else with a sensible string form.

```
namo[symbol: /^BH/]          # symbols starting with BH
namo[symbol: /gold/i]        # case-insensitive match
namo[sector: /mining|resources/i] # alternatives
namo[symbol: /^BH/, sector: 'Energy'] # regex + exact, composable
```

Implementation is a single `when` `Regexp` branch in `Row#match?`:

```
when Regexp
  coordinate.match?(self[dimension].to_s)
```

Same weight as adding `proc` support — one additional `when` branch in the same case statement.

The `to_s` coercion has predictable behaviour across types:

value	.to_s	matches //	matches /. /
<code>nil</code>	<code>""</code>	yes	no
<code>42</code>	<code>"42"</code>	yes	yes
<code>10.0</code>	<code>"10.0"</code>	yes	yes
<code>:priority</code>	<code>"priority"</code>	yes	yes
<code>Date.new(...)</code>	<code>"2026-05-21"</code>	yes	yes

Regex is more ergonomic than the equivalent `proc` for string pattern matching:

```
# Regex
namo[symbol: /^BH/]

# Equivalent proc
namo[symbol: ->(v){v.to_s =~ /^BH/}]
```

The regex form is shorter, more declarative, and immediately legible. It doesn't replace procs — procs handle arbitrary logic — but for pattern matching on string-valued (or string-coercible) dimensions it's the natural tool.

Regex composes with all other selection types in the same `[]` call: exact values, arrays, ranges, procs, projection, and contraction.

0.9.0 (2026-05-21): Composition operators (`*`, `**`, `/`)

The dimensional composition algebra. Three operators that extend `Namo` from the same-dimensions algebra (set operators in 0.4.0–0.5.0, comparison operators in 0.6.0) to combining and decomposing `Namos` with different dimensions.

*** (equi-join on shared dimensions)** Pairs rows where coordinates match on every shared data dimension. Inner-join semantics — unmatched rows from both sides are dropped. Output dimensions are `self.data_dimensions` followed by other's exclusive dimensions; output multiplicity is the product of input multiplicities on each matching key.

```
ohlcv * fundamentals # joins on shared :symbol
```

Requires at least one shared data dimension. No overlap raises `ArgumentError` — silently falling through to a Cartesian product would turn a logic error into a large pile of nonsense rows. Formulae merge from both sides; self wins on conflict.

**** (Cartesian product)** Every row from the left paired with every row from the right. Output has `self.data.length * other.data.length` rows; output dimensions are `self.data_dimensions + other.data_dimensions`.

```
products ** quarters
```

Requires **no** shared data dimensions — the precondition is the mirror image of `*`. Any overlap raises `ArgumentError`. The visual relationship is deliberate: `*` is the filtered version, `**` is the explosive version — more sigil, more output. Formulae merge from both sides; self wins on conflict.

/ (decomposition) Removes from self the dimensions that are also in other (the intersection), then dedupes the projected rows. The inverse of `*` and `**`.

```
combined / fundamentals # removes shared dimensions, keeps everything else
```

No precondition — `/` is total on `Namo × Namu`. When the operands share no dimensions, the intersection is empty and `self / other` returns a `Namo` equal to self. The asymmetric strictness — `*` and `**` raise, `/` is loose — reflects a structural distinction: `*` and `**` are *combining* operators that need a specific relationship between operands to produce a meaningful result, while `/` is a *projecting* operator where “project away nothing” has a natural answer (“return the original”).

The looseness earns `/` algebraic properties a strict version would lose: identity test (`c / b == c` iff they share no dimensions), idempotence (`(c / b) / b == c / b`), and pipeline composition (a step that applies `/` separator can run over any `Namo` without special-casing applicability).

Round-trip identity:

- `(a ** b) / b == a` exactly.
- `(a * b) / b == a[-:shared]` — the dimensions shared with `b` are lost on decomposition.

The asymmetry between the two round-trip cases is real: `/` operates only on the two values it receives and cannot distinguish “shared dimension that belonged to both” from “exclusive dimension that belonged only to the right”. Removing the intersection is the only rule expressible from the operands alone.

0.10.0 (2026-05-28): Row comparison

Extends 0.6.0’s comparison work one level down. 0.6.0 settled comparison at the `Namo` level — `==`, `eq1?`, `hash`, `===`, and the subset/superset operators — but left `Row` without value semantics. The omission was defensible at the time: `Row` read as implementation detail behind `Namo`’s algebra. Applied use showed otherwise. Rows leak through each, get reached for directly in interactive sessions, and become prerequisites for the dedup, hash-keying, and Row-against-Row equality that 0.11.0’s `Enumerable` coherence pass needs. This release closes that gap: `Row` gets `==`, `eq1?`, and `hash` matching its role as a hash-shaped value.

The pattern of 0.10.0 revisiting 0.6.0 is worth flagging because 0.11.0 revisits 0.2.0 in the same way — the substrate beneath an earlier algebra needs to catch up to applied use. The “Notes on these two releases” at the end of 0.11.0 expands on the pattern.

== Data equality. Two Rows are equal if their underlying `@row` hashes are equal. Formulae are not part of equality — they’re attached by the surrounding `Namo`, not properties of the row.

```
def ==(other)
  other.is_a?(Row) && @row == other.to_h
end
```

This mirrors `Namo#==` one level down. 0.6.0's `Namo#==` ignores class and formulae and compares data; `Row#==` follows the same rule. Two Rows with identical `@row` data are equal regardless of which `Namo` they came from.

eql? Mirrors `Row#==` in shape — an `is_a?(Row)` gate followed by a hash comparison. Unlike `Namo#eql?`, there's no class-identity gate (no `self.class == other.class`); Rows don't have a subclass hierarchy with included modules, so there's no class story to enforce.

```
def eql?(other)
  other.is_a?(Row) && @row.eql?(other.to_h)
end
```

The difference from `Row#==` is that `eql?` uses `Hash#eql?` for the underlying comparison rather than `Hash#==`, matching Ruby's convention (`1 == 1.0` is true but `1.eql?(1.0)` is false; Hash comparison follows from its keys' and values' comparison). The Row-level consequence is that `Row.new({n: 1}, {})` `==` `Row.new({n: 1.0}, {})` is true but `.eql?` between the same two is false.

hash Consistent with `eql?` — Rows that are `eql?` produce the same hash, making them usable as Hash keys and Set members.

```
def hash
  @row.hash
end
```

Why these three, not the full Nammo stack `Namo` gets `==`, `eql?`, `hash`, `===`, `<`, `<=`, `>`, `>=` because a `Namo` is a collection of rows with set-theoretic relationships. A `Row` is a record — a single hash-shaped value. The set-theoretic operators don't translate: a `Row` isn't a "subset" of another `Row` in any meaningful sense, and `===` for pattern-match dispatch doesn't apply to row-level values the way it applies to analytical shapes. The three that translate are the value-semantics trio: `==`, `eql?`, `hash`. Those are what a hash-shaped value needs to participate in Ruby's collection machinery correctly.

0.11.0 (2026-05-31): Enumerable methods return Namos

Extends 0.2.0's `Enumerable` inclusion. 0.2.0 made `Namo` `Enumerable` but left Ruby's default return types in place — `select`, `reject`, and `friends` produced Arrays, breaking the analytical chain. Applied use made the friction visible: every interactive session that selects rows from a `Namo` and tries to continue working with the result hit a `NoMethodError` for selection or projection on the Array. 0.11.0 specialises the subset-returning `Enumerable` methods to wrap in `Namo`. The Row-equality groundwork for `uniq` and `partition` came in from 0.10.0.

These are sequence-view operations: `select`, `reject`, `sort_by`, `first(n)`, `last(n)`, `take(n)`, `drop(n)`, `take_while`, `drop_while`, `uniq`, and `partition` all care about row order and produce ordered subsets. They sit alongside the set-view operators (`==`, `<`, `&`, `|`, etc.) from 0.4.0–0.6.0. `Namo`'s dual nature — set when membership is what matters, sequence when order is what matters — is now realised across both families: set operators ignore order and produce set-correct results; `Enumerable` methods respect order and produce ordered results. The same `Namo` supports both views.

```
# Before (0.2.0–0.10.0)
filtered = namo.select{|row| row[:close] > 40.0}
filtered.class      # => Array
filtered[:symbol: 'BHP'] # => NoMethodError

# After (0.11.0)
filtered = namo.select{|row| row[:close] > 40.0}
filtered.class      # => Nammo
filtered[:symbol: 'BHP'] # works – selection, projection, formulae, everything
```

Scope Methods that return Namos:

- `select`, `reject` — predicate-based subset, formulae carry through. `select`'s Enumerable aliases `filter` and `find_all` are aliased to the override, so they return Namos too — the only overridden method that has aliases to keep in step.
- `sort_by` — reordered Namo, formulae carry through.
- `first(n)`, `last(n)` — leading and trailing subsets. Without an argument, return a single Row (or nil for empty); with an argument, return a Namo.
- `take(n)`, `drop(n)` — leading subset and its complement.
- `take_while`, `drop_while` — predicate-based leading subset and its complement.
- `uniq` — dedupe rows on full-row equality (`Row#==/ eql?` from 0.10.0). With a block, dedupe on the block's return value, following Enumerable#`uniq`'s convention.
- `partition` — returns `[Namo, Namo]` — matches and non-matches.

Methods that did not change: `map`, `flat_map` (transformed values, possibly not row-shaped), `reduce`, `sum`, `min_by`, `max_by`, `count` (scalars), and `each` (already returns an Enumerator or yields Rows).

`group_by` is deliberately excluded, and not merely deferred — it is *structurally blocked*. Every method above returns a Namo (or `[Namo, Namo]`), and Namo exists. `group_by` returns a keyed set of sub-Namos, and the right type for that is `Namo::Collection`, which doesn't exist until 0.18.0. It lands at 0.20.0.

Implementation pattern Subset-returning methods construct via `self.class.new` so subclass type is preserved; formulae carry through, duped to avoid shared mutation.

```
def select(&block)
  self.class.new(@data.select{|row| block.call(Row.new(row, @formulae))}, formulae: @formulae.dup)
end

def first(n = nil)
  if n
    self.class.new(@data.first(n), formulae: @formulae.dup)
  else
    @data.first ? Row.new(@data.first, @formulae) : nil
  end
end

def uniq(&block)
  rows = block ? @data.uniq{|row| block.call(Row.new(row, @formulae))} : @data.uniq
  self.class.new(rows, formulae: @formulae.dup)
end

def partition(&block)
  matches, non_matches = @data.partition{|row| block.call(Row.new(row, @formulae))}
  [
    self.class.new(matches, formulae: @formulae.dup),
    self.class.new(non_matches, formulae: @formulae.dup),
  ]
end
```

The no-block `uniq` dedupes raw `@data` hashes via `Array#uniq`, which uses `eql?/hash` — exactly what `Row#eql?/Row#hash` delegate to — so it matches Row equality while avoiding Row allocations. The practical consequence is numeric type strictness: `Namo.new([{:n: 1}, {:n: 1.0}]).uniq` keeps both rows, matching `Row#eql?`.

As of 0.11.1, these methods (and `each`) live in a `Namo::Enumerable` module included into Namo, rather than inline on the class — a behaviour-preserving reorganisation. The module include `::Enumerables` itself

and sits above `stdlib Enumerable` in Namo's ancestor chain, so the overrides win while `map/reduce/etc.` fall through unchanged.

Subclass considerations Same trap as the operators — `self.class.new(...)` constructs an instance of the subclass, firing its `initialize` side effects. The `if name guard` pattern (introduced in 0.12.0 alongside the `name: attribute`) is the documented convention for subclasses with side effects in `initialize`.

Backward compatibility A breaking change for code expecting `Array` returns from these methods. In 0.x, breaking changes are allowed; the release note: “Subset `Enumerable` methods now return `Namos`. If your code expected `Arrays`, call `.entries` or `.to_a` on the result.”

Performance note `last(n)` goes straight to `@data.last(n)`, the efficient path. No fall-through to `Enumerable`'s materialise-then-slice behaviour.

Notes on these two releases 0.10.0 and 0.11.0 are a single conceptual move. 0.10.0 extends 0.6.0's comparison work to `Row`; 0.11.0 extends 0.2.0's `Enumerable` inclusion to make subset-returning methods produce `Namos`. They pair because 0.11.0's `uniq` and `partition` need `Row` equality to work correctly, and 0.10.0 puts that in place.

The pattern of revisiting earlier releases is worth flagging because it recurs. Namo's earlier releases closed algebras cleanly at one conceptual level — `Enumerable` at 0.2.0, the set algebra at 0.4.0–0.5.0, comparison at 0.6.0, composition at 0.9.0. Each was a coherence statement: a complete, deliberately-bounded piece of algebra. Theoretical design closes algebras at their natural level; the substrate beneath the algebra (return types, value semantics on the underlying objects) tends to get left at Ruby defaults because the design effort goes into the structural completeness above.

Applied use stresses different parts. Interactive sessions reveal that `select` returning `Array` breaks chaining. Wanting to dedupe rows reveals that `Row` never got the value semantics the Namo level got. The substrate beneath the algebra needs to catch up to applied use. This is healthy: the earlier releases are sound — the revisits are extensions, not corrections. If applied use were revealing genuine design errors that would be different; instead it reveals where deliberate scope limits in earlier releases were tighter than the applied surface needs. Future releases may follow the same pattern.

0.12.0 (2026-06-01): Constructor widening — keyword `data` and `name`:

One coherent edit to `initialize`: the constructor grew two optional keyword arguments. `Data` can now be passed by the `data:` keyword as well as positionally, and a `Namo` can carry a `name`. Both are additive — every existing call site (`Namo.new([ {...} ])`, `Namo.new([ {...} ], formulae: {})`, the no-arg and formulae-only forms, and the operators' internal `self.class.new(rows, formulae: ...)`) is unaffected.

```
# Before (through 0.11.x)
def initialize(data = [], formulae: {})

# 0.12.0 – widened
def initialize(positional_data = nil, data: [], formulae: {}, name: nil)
  @data = positional_data || data
  @formulae = formulae
  @name = name
end
```

Positional data wins when both positional and keyword `data:` are present. The positional default changed from `[]` to `nil` so it acts as a “not supplied” sentinel that lets `|| data` fall through to the keyword path; an explicit `Namo.new([])` still yields `@data == []` because the truthy empty array short-circuits.

```
Namo.new([{:x: 1}])           # positional – unchanged
```

```
Namo.new(data: [{x: 1}])      # keyword – new
Namo.new([{:x: 1}], data: [{x: 2}])  # positional wins → @data == [{x: 1}]
```

name: is the load-bearing half — the driver is `Namo::Collection` (0.18.0), where members identify themselves so the collection can find them by name, replace them on re-add, and label them in summaries. `Namo` gains `attr_accessor :name`; `name=` gives post-construction mutation. Operator-derived `Namos` (`+`, `*`, `select`, ...) construct without `name`, so their `@name` is `nil` — a derived object is not the original, and giving it the parent's name would mislead.

That nil-on-derivation behaviour is what makes the subclass guard convention work: subclasses guard initialize side effects with `return unless name`, so operator-derived instances skip them and only explicitly-named constructions fire them. `super` with no parentheses forwards every argument unchanged.

```
class TradingAnalysis < Namu
  def initialize(positional_data = nil, data: [], formulae: {}, name: nil)
    super
    return unless name
    register_indicators
  end
end
```

This is the answer to the “operator return type” open question for subclasses with side effects in `initialize` — they need not override every operator to stop the result of `*` or `select` from re-running construction; they guard on `name`. The convention is documented in the README and applied throughout `Namo`'s own subclasses.

These two changes shipped together because they are the same modification: the constructor's signature growing optional keyword parameters. Polymorphic `[]=` (0.13.0) was deliberately not bundled — it changes assignment dispatch, a different method and a different concern. The seam is “constructor-signature changes together; assignment-dispatch changes separately.”

0.13.0 (2026-06-01): Polymorphic `[]=`

`[]=` dispatches on the type of the value assigned. A `Proc` registers a formula; anything else broadcasts the value to every row. The two branches are mirror images that together enforce **exclusive storage** — a name is either a data dimension or a derived dimension, never both.

```
def []=(name, value)
  case value
  when Proc
    @data.each{|row| row.delete(name)} if @data.first&.key?(name)
    @formulae[name] = value
  else
    @formulae.delete(name)
    @data.each{|row| row[name] = value}
  end
end
```

The `Proc` branch clears any data column of that name before registering the formula; the `else` branch deletes any formula of that name before broadcasting. Last write wins, and there is no shadowing. The `Proc`-branch clear is guarded (`if @data.first&.key?(name)`) so it walks the rows only when there's actually a data column to clear; the `&.` also handles the empty-data case where `@data.first` is `nil`. The `else` branch's `@formulae.delete(name)` and broadcast are unconditional — `Hash#delete` on an absent key is a cheap no-op, and the broadcast is the operation's actual work. The match is on `Proc` specifically (`Proc === value`), consistent with `Row#match?`'s `proc-selection` branch from 0.8.0, so an array is not a `proc` and `namo[:weights] = [1, 2, 3]` broadcasts the array as each row's value.

Exclusivity is self-enforcing because the inspection vocabulary is derived live. `data_dimensions` reads the

keys of the first row and `derived_dimensions` reads the keys of `@formulae`, so clearing the data column is what removes a name from `data_dimensions`, and deleting the formula is what removes it from `derived_dimensions`. There's no separate bookkeeping to keep in sync. A name assigned a scalar shows up in `data_dimensions`; a name assigned a proc shows up in `derived_dimensions`; never both, so it appears in `dimensions` (their concatenation) exactly once.

Why polymorphic, not a separate method The alternative — `[]=` for formulae plus a separate `broadcast/set_all` for scalars — was considered and rejected. `[]` already dispatches polymorphically over exact values, arrays, ranges, procs, and regexes (0.8.0); `[]=` follows the same convention rather than introduce a parallel naming scheme. The exclusivity rule then lives at the one place assignment happens, instead of being pushed onto the user, who could otherwise leave a formula shadowed by a data column of the same name (or vice versa) — stale state that, with the live-derived inspection methods, would surface as a name listed in both `data_dimensions` and `derived_dimensions`. And it reads naturally: `namo[:status] = 'active'` is “set status to active across this Namu,” which is exactly what it does.

This shipped separately from the 0.12.0 constructor widening by design: that release changed the constructor's signature; this one changes assignment dispatch — a different method and a different concern. The seam is “constructor-signature changes together; assignment-dispatch changes separately.”

0.14.0 (2026-06-08): Blocks on composition operators (*, **)

`*` and `**` take an optional block that refines which rows pair. Without a block, both behave exactly as they did in 0.9.0 — `*` pairs every shared-dimension match, `**` pairs every row with every row. With a block, the operator hands it the current left row and a Namu of candidate right rows, and the block returns the subset to pair.

```
block.call(row, candidates) # => Namu
```

`row` is the Row for the current left row, carrying self's formulae (the same object each yields), so `row[:date]` and any self formula resolve inside the block. `candidates` is a Namu of right rows carrying other's formulae — for `*`, the rows already matched on the shared data dimensions; for `**`, all of other's rows, with no pre-filter — so the block can select on other's derived dimensions as readily as its data. The block returns a Namu of the rows to pair, each merged into the left row. It is a selector, not a reducer: zero, one, or many rows are all valid returns. An empty returned Namu pairs nothing, dropping that left row — inner-join semantics preserved.

The canonical use is a matching rule plain `*` cannot express because it pairs every match: match each daily price to a single quarterly report — the most recent one dated on or before it.

```
prices.*(quarterly) do |row, candidates|
  candidates[quarter_end: ->(qe){qe <= row[:date]}].sort_by{|f| f[:quarter_end]}.last(1)
end
```

`**`'s block is the parallel for the no-shared-dimensions case — a conditional product, pairing each order with only the shipping tiers that can carry it:

```
orders.**(tiers) do |row, candidates|
  candidates[max_weight: ->(w){w >= row[:weight]}]
end
```

The block decides only which rows pair; it does not touch the result's formulae. Those are set once by the operator, `other.formulae.merge(@formulae)`, identical to the no-block path — `candidates` carries other's formulae only so block selection on other's derived dimensions resolves, read-only and local to the block, and the returned Namu contributes row data only. So formula handling is byte-identical between the block and no-block forms. A derived dimension of other (a formula, not a stored key) does not become stored data on merge; it reappears on the result because the result carries other's formulae, resolving to the same value. The operators' preconditions are unchanged and additive: `*` still requires a shared dimension and `**` still requires disjoint dimensions, block or no block. A `**` block that returns its candidates unchanged reproduces

the no-block product, and one that matches on shared dimensions by hand reproduces no-block * — * is ** with the shared-dimension match applied first.

The governing principle A block form is warranted exactly when the operation gives consideration to a dimension in isolation. The composition operators do — * singles out the shared dimensions as the join axis and the block reasons about particular dimensions’ values; ** gives no such consideration itself, and the block is where the user supplies it. The set operators (+, -, &, |, ^) do not: they act on the whole row as an indivisible value (whole-row equality via eql?/hash), singling out no dimension, so there is nothing for a block to refine or supply. Decomposition (/) considers dimensions collectively — a set of names operated on wholesale — never a dimension in isolation, so it too takes no block.

This ties to orthogonality: where a conditional operation can be expressed by composing existing orthogonal operations, it should be, rather than by adding a block to an operator. The anti-join a set-operator block would have served is one such case — `today[symbol: ->(s){!excluded.include?(s)}]` (with `excluded` a Set for an `O(today + exclusions)` lookup) expresses it using only shipped features, and is at least as efficient as a block form would be. A block earns its place only where no such composition exists — precisely the composition operators, where the dimension-in-isolation matching is underdetermined and cannot be assembled from existing parts. The comparison operators are excluded on the same ground: a keyed comparison is already `today[:symbol] == yesterday[:symbol]`, a comparison of projections, so no block is needed; and `eql?` would not take one regardless, its job being the strictest equality.

0.15.0 (2026-06-12): Two-arity formulae

Row carries a reference to the Namo that yielded it, and `Row#[]` dispatches on a formula’s arity. A proc with arity 1 (and arity 0, and any negative arity) is called `proc.call(row)` exactly as before — row-scoped. A proc with arity **exactly 2** is called `proc.call(row, namo)`, where `namo` is the Namo the row belongs to — collection-scoped, so the formula can reach the rest of the dataset and compute across rows: moving windows, ranks, running totals.

```
prices[:sma] = proc do |row, namo|
  window = namo[symbol: row[:symbol], date: ->(d){d <= row[:date]}]
  window.values(:close).sum / window.count.to_f
end
```

`namo` is the yielding Namo, live — every consequence follows from “whichever Namo constructed this Row is the parent”. A filtered Namo’s rows window over the filtered rows; an operator result’s rows window over the result; appending a row changes every two-arity value on the next access, with no caching, per the live-computation discipline. A two-arity formula that references its own name recurses unguarded, exactly as a self-referential one-arity formula already does.

The dispatch pins exactly arity 2. Negative arities — `->(row, namo = nil){}` and `proc{|row, *rest|}` are both arity -2 — take the one-arity path in this release; whether a trailing splat or optional should be collection-scoped is settled in 0.17.0’s arity > 2 generalisation, not pre-empted here. The case on arity is the seam that release extends.

Row’s constructor gains an optional third parameter, `namo`, defaulting to `nil`, so every existing `Row.new(row, formulae)` call site keeps working. Enumeration and predicate evaluation (`each`, `select` and its aliases, `reject`, `sort_by`, `first`, `last`, `take_while`, `drop_while`, `uniq`’s block form, `partition`) pass `self` as the parent, as do `values_for`’s derived branch and the `*/**` block paths — so `values`, `coordinates`, `to_h`, `selection`, and composition blocks all resolve two-arity dimensions. A Row constructed without a Namo raises a clear `ArgumentError` naming the formula when a two-arity formula is asked of it, rather than letting `nil` leak into the formula body.

A two-arity formula’s body typically scans the parent, so materialising a full column is $O(n^2)$ -shaped. This is the accepted pure-live cost; benchmarking (1.1) measures it and freeze-gated caching (2.x) relieves it. No memoisation in this release.

0.16.0 (2026-06-12): Data/formula exclusivity

A name is data or derived, never both. []= already enforced this at assignment — a proc deletes the data column of the same name, a scalar deletes the formula. Two library paths violated it, and the violation surfaced as a precedence disagreement between the access paths: `values_for` resolves a name data-first, `Row#[]` formula-first, so an aliased `Namo` answered differently depending on how it was asked. This release extends enforcement to both paths.

Projection. Through 0.15.0, `prices[:date, :sma]` materialised the formula’s values into the projected rows and carried the formula forward. The result answered `values(:sma)` from the stored data but recomputed `first[:sma]` through the formula — over the projected `Namo`, whose input columns (`:symbol`, `:close`) were just dropped. Row access and selection on the projected dimension raised.

The fix: projection drops the formulae it materialises.

```
carried = positive.any? ? @formulae.reject{|name, _| positive.include?(name)} : @formulae.dup
self.class.new(projected, formulae: carried)
```

All access paths then agree, `dimensions` lists the name once, and a dependent formula *not* named in the projection carries through and resolves off the materialised column — `project[:date, :sma]` with a `:double_sma` formula referencing `:sma`, and `:double_sma` keeps working against the stored values. No liveness is lost: projection already snapshots data into fresh hashes (unlike selection, which shares row objects with its source), so “values current at the moment of projection” was its semantic for data dimensions all along; derived dimensions now follow the same reading. Materialisation goes through the Rows the source yields, so a two-arity formula materialises windowed over the yielding `Namo`, per 0.15.0’s discipline — and with selections in the same call (`prices[:date, :sma, date: 2..3]`), the yielding `Namo` is the filtered one, so the values window over the selection, consistent with select-then-project. Only the positive-projection branch changed: contraction and selection-only calls carry all formulae unchanged.

The rule is also a control surface: materialisation is selective, and the projection list is the selector. Naming a derived dimension asks for its values — a stored snapshot, computed against the source. Omitting it leaves it as computation — formulae not named in the projection carry through live and compute from the projected columns on every access, exactly as they carry through selection. `sales[:price, :quantity, :revenue]` returns a `Namo` with `:revenue` as stored values; `sales[:price, :quantity]` returns one where the `:revenue` formula recomputes from the result’s own rows whenever asked. All four combinations of materialised/live × inputs-present/inputs-dropped are expressible by what is named, and the only failing one — carrying a formula whose inputs the projection cut — is the user’s explicit choice, the same caveat-emptor as contracting away a formula’s inputs. The `:double_sma` case above is the boundary worth knowing: a carried formula that references a materialised dimension is half-frozen — itself live, computing over snapshot inputs — which follows directly from the rule. An explicit carry-live marker (a `~:revenue` wrapper mirroring `~:dim` contraction) was considered and rejected as redundant: “name the inputs, omit the formula” already spells it, and duplicate spellings cut against orthogonality.

Composition. `*` and `**` merge formulae from both sides, so one operand’s data dimension could collide with the other’s derived dimension — `:margin` as an audited stored figure on the left, `:margin` as a formula over `:cost` and `:price` on the right. The result held both, silently: `values(:margin)` gave the stored figure, `first[:margin]` the computed one.

The fix: `*` and `**` raise `ArgumentError`, naming the colliding dimensions, when `(data_dimensions & other.derived_dimensions) | (derived_dimensions & other.data_dimensions)` is non-empty — block and no-block forms alike, since preconditions are block-independent, per the 0.14.0 rationale. The operands disagree about what the name means — stored fact on one side, computation on the other — and there is no last-write order to appeal to, so this is precisely where the operators’ existing character applies: refuse the ambiguous operand pair loudly rather than guess. The user resolves the collision explicitly before composing, by contraction (`audited[-:margin] * modelled`) or projection.

The other operators need nothing. The set operators’ matching-data-dimensions precondition already blocks the asymmetric case — if one side carries the name as data and the other doesn’t, their data dimensions can’t

match. Formula-vs-formula collisions stay left-wins, as documented — that’s resolution, not aliasing. The constructor stays unguarded: `Namo.new(data, formulae: ...)` can still hand-build an aliased `Namo`, the same trust it already extends to row shapes.

Rejected alternatives. Data-first precedence in `Row#[]` (matching `values_for`) fixes the symptom everywhere in one line, but legitimises the aliased state instead of eliminating it: formulae on hand-built `Namos` go silently dead behind same-named data, and `dimensions` keeps listing the name twice. Keeping projection live — carrying the formula, not materialising — fails structurally: the formula’s inputs are exactly what projection dropped, and with opaque procs there is no inferring whether they survived (the same wall as proc comparison in `===`).

That last point is the one to revisit. If dependencies were knowable, projection could keep a formula live when its inputs survive the cut and materialise only when they don’t. The 2.x bare-names DSL makes dependencies observable by construction rather than inferred, so dependency-aware live projection has a natural home there.

Parameterised formulae (0.17.0) cannot be materialised without their arguments; what projection does with an arity > 2 dimension is defined in that release, extending this rule.

0.17.0 (2026-06-13): Parameterised formulae

A formula can declare parameters beyond `(row, namo)`, and the arguments arrive at access time: `Row#[]` grows a trailing splat and forwards what it’s given. One definition serves every column and every setting, completing the formula trajectory — single-row (0.1.0), cross-row (0.15.0), parameterised here — through the same mechanism, with the same currentness.

```
prices[:sma] = proc do |row, namo, field, period|
  window = namo[symbol: row[:symbol], date: ->(d){d <= row[:date]}].last(period)
  window.sum{|r| r[field]} / window.count.to_f
end
```

```
prices.last[:sma, :close, 20]    # Row inserts self and namo, forwards :close and 20
prices.last[:sma, :volume, 50]   # the same formula, different field and period
```

The dispatch rule: required parameters decide scope 0.15.0 pinned dispatch on exactly arity 2 and deferred the negative arities to this release. The generalisation that settles them: **the number of required parameters decides a formula’s scope**. One (or none) means row-scoped, called `formula.call(row)`. Two or more means collection-scoped, called `formula.call(row, namo, *arguments)` — everything past the second parameter receives the call-site arguments. For a proc of arity $n \geq 0$ the required count is n ; for negative arity it is $-n - 1$, Ruby’s encoding of “required parameters before the splat or optionals”.

The principle: what a formula *requires* is the only signal the proc object carries about what it means to receive. Optionals and splats express openness, not demand, so they can’t claim a calling convention — `->(row, namo = nil){}` and `proc{|row, *rest|}` (both arity -2, one required parameter) stay row-scoped, exactly as 0.15.0 left them, while `proc{|row, namo, *fields|}` (arity -3, two required) is collection-scoped with optional arguments. The 0.15.0 case on exact arity becomes a predicate on required count; no shipped behaviour changes, the pinned-open cases close.

Argument-count enforcement at `Row#[]` Every dimension now has an argument signature, and `Row#[]` enforces it: data dimensions and row-scoped formulae take none; a fixed-arity collection-scoped formula takes exactly arity $- 2$; a splatted one takes at least its required count minus two. The wrong number raises an `ArgumentError` in Ruby’s own phrasing — wrong number of arguments for `:sma` (given 0, expected 2), (given 1, expected 0) for arguments handed to a data dimension, (given 0, expected 1+) for an underfilled splat.

Enforcement is needed because Ruby wouldn’t provide it: non-lambda procs silently fill missing parameters with `nil` and discard extras, so an underfilled parameterised formula would compute against `nil`s — the same

silently-two-answers class of bug whose data/formula form 0.16.0 eliminated. The check lives at `Row#[]` because that is where names are resolved, the same reasoning that put exclusivity at `[]=`. One guard covers every access path — direct row access, selection predicates, materialisation — because they all pass through it.

The Namu-context guard from 0.15.0 generalises alongside: any collection-scoped formula (two-arity or parameterised) asked of a Row constructed without a Namu raises the renamed message, `collection-scoped formula :sma requires a Namu context`, but this Row has none. Argument count is checked before context — a call of the wrong shape is wrong regardless of where the Row came from.

Materialisation: the 0.16.0 rule extended 0.16.0 left one question open: a formula that requires arguments cannot be materialised without them, so what do projection and the bulk views do with one? The rule splits on whether the ask is explicit:

- **Asking by name raises.** `values(:sma)`, `coordinates(:sma)`, `naming :sma` in a projection, and selecting on `sma`: all raise the argument-count `ArgumentError` — there is no argument channel in any of those calls, and inventing one was rejected (below). The raise comes from `Row#[]` naturally; no second enforcement site.
- **The bulk views omit.** No-argument `values`, `coordinates`, and `to_h` return every dimension that can be materialised and skip those that can't — the honest subset, not an error, because registering one parameterised formula shouldn't poison whole-Namu views that are correct for every other dimension. `dimensions` and `derived_dimensions` still list the name: those describe the queryable namespace, and the dimension is queryable — with arguments.

The boundary is *requires*, not *accepts*: a splat directly after `namu` (`proc{|row, namu, *rest|}`) requires nothing extra, so it materialises with no arguments and appears in the bulk views, while `proc{|row, namu, field, *rest|}` requires one and is omitted. Internally the bulk views select via a private `materialisable_dimensions`; promoting it to the public inspection vocabulary waits for a use case, per the 0.7.0 precedent of not shipping accessors ahead of need. The empty-Namu case stays lazily honest, returning `[]` without invoking the formula, as 0.15.0 established for two-arity.

To materialise particular values, bind the arguments in a row-scoped wrapper and name that — `prices[:sma_close_20] = proc{|row| row[:sma, :close, 20]}` re-enters the materialisable world, and 0.16.0's projection control surface applies to the wrapper unchanged. This composition is why no argument channel was added to `values` or projection (a `values(:sma, arguments: [...])` form was considered): the wrapper spells the binding with shipped parts, names the bound result so it can be projected, selected, and carried like any other formula, and keeps one calling convention for `values` rather than two.

A row-scoped formula can call a parameterised one with arguments — the wrapper above is the canonical case — and a parameterised formula can reference any other formula through `row`. Self-reference recurses unguarded, as everywhere in the formula mechanism.

0.18.0 (2026-06-13): Namu::Collection

`Namu::Collection < Namu` is the first member of the family beyond `Namu` itself: a hierarchical aggregate that holds an Array of named member `Namos` (`members`) and exposes summary and detail views across them. It lives at `lib/Namu/Collection.rb` per file convention, and it is the return type that unblocks `group_by` at 0.20.0.

```
class Car < Namu::Collection
  def summary(dimension, by: :assembly, reducer: :sum)
    super
  end

  def detail(by: :assembly)
    super
  end
end
```

```

end
end

class SubAssembly < Namo; end

gt = Car.new
gt << [powertrain, chassis, body]           # each a named SubAssembly
gt.summary(:weight).values(:weight).sum     # total by summing assembly summaries
gt.detail.values(:weight).sum               # total by summing every line item

```

Members are the substance; the data view is lazy detail A Collection’s substance is members; the inherited `@data` is a derived view of them, not independently-stored state. The lazy view is **detail** — the lossless union of the members’ rows — because a Collection’s rows simply *are* its members’ rows, while a summary is a reduction posed against them and so is never reached by accident (only through `summary/as_summary`). Any inherited row-operation (selection, projection, `each`, `values`, the operators) reads the view, so `collection[component: 'engine']` just works, behaving as the detail.

`<<` accepts a single member or an array (via `flatten`). A member whose name collides with an existing one replaces it — last-write-wins, making the name $\rightarrow$ member mapping a dictionary, not a multimap. `find(name)` returns the member with that name or `nil`.

The four views, and the inject-iff-absent rule `summary(dimension, by:, reducer:)` and `detail(by:)` are non-mutating, returning a fresh Namo derived from the members; `as_summary(dimension, by:, reducer:)` and `as_detail(by)` are mutating, setting the Collection’s data to the chosen view and returning self for a fluent step. `summary` reduces each member to a `{by => member.name, dimension => reduced}` row; `reducer:` is any method the column responds to (`:sum` default, `:mean` via a statistics gem). `detail` unions the members’ rows under the single conditional `row.key?(by) ? row : row.merge(by => member.name)` — intrinsic dimensions pass through, an absent label is injected from the member name. This is the one place the assembly path (`<<`, extrinsic names) and the partition path (`group_by`, intrinsic names — 0.20.0) meet, and it makes `as_detail(:assembly)` the dimension-creating step that promotes a member name into retained data; removal is only ever an explicit contraction (`[-:assembly]`).

Names enforced at use, not insertion `<<` has no insertion guard against unnamed members. An unnamed member is appended (no name to collide on) and is simply unfindable by `find` — the honest consequence of having no name, not an error. This is deliberate: `group_by` (0.20.0) will construct members named by their group value, and a `nil`-valued group key produces a legitimately `nil`-named member whose rows still materialise and round-trip via the intrinsic grouping dimension. Forbidding `nil` names eagerly would break that case; moving enforcement to use-site avoids it, matching Namo’s compute-on-access discipline.

Resolved at implementation Two points the design left mechanically under-determined against the shipped source were settled here:

- **Materialisation model — rebuild-on-`<<`, not recompute-per-access.** The inherited row-operations read the `@data` ivar directly, so there is no transparent per-access seam without rerouting every shipped operator through the data accessor. Rather than touch the 0.2.0–0.17.0 internals for a refactor whose only gain was literal wording, `<<` (the sole members mutator) rebuilds `@data = detail.data`. Because it is the only mutator, this is observationally identical to per-access recomputation across the documented surface — and it lets a mutating `as_*` view mean what it reads like. The consequence, diverging from the originally-drafted “pure-live per-access / transient `as_*`” wording: an `as_summary/as_detail` view **persists until the next `<<`** (which re-materialises detail), rather than being discarded by the next row-operation. `as_summary(:weight).values(:weight)` therefore reads the summary, as intended. Freeze-gated memoisation remains a 2.x optimisation, now attached to a settled rebuild-on-mutation view.

- **find shadows Enumerable#find on Collections.** `find(name)` is member lookup; a `Collection`'s natural elements are its members, and the call reads well (`gt.find(:powertrain)`). It shadows the inherited `Enumerable#find` (alias `detect`) on `Collections` only — `detect` survives for predicate search, and `find` with a block fails loudly on arity rather than misbehaving. `as_detail`'s label argument is positional (`as_detail(:assembly)`), as it carries no dimension; `summary/detail/as_summary` keep by: as a keyword alongside their other arguments. (Superseded at 0.28.0: `detail` and `as_detail` both take the label positionally or by keyword, per the constructor precedent — “no dimension” didn’t distinguish the twins.)

0.19.0 (2026-06-14): Row-multiset equality

The comparison operators were defined as multiset-theoretic on rows at 0.6.0, but the *mechanism* was a sorted canonical form: `canonical_data` sorted `@data` by `row.values_at(*data_dimensions.sort)`, and `==`, `eql?`, `hash`, and the subset/superset operators (`<`, `<=`, `>`, `>=`) compared those sorted sequences. A sort needs a total order over values, and `nil` (like `NaN`) has no `<=>` against a present value — so the moment a dimension held a mix of `nil` and present values the sort raised `ArgumentError`, and two such `Namos` could not be compared at all.

A multiset needs no ordering. `canonical_data` is replaced by `row_multiset` (`@data.tally`), and the same operators now compare row multisets keyed by `Hash#hash/#eql?`. This reaches the already-decided multiset semantics directly rather than via a sort, so the `nil/NaN` crash is gone — not patched, but structurally absent, because nothing is ordered. Equality is row-order-blind and dimension-order-blind for free (`hash` equality ignores key order), exactly as the comparison operators were always meant to be.

It also unifies the relation. The subset operators already tallied (keyed on `eql?`), while `==/eql?/hash` rode `Hash#==` through the sorted array — so the family disagreed on numeric type: `{a: 1} == {a: 1.0}` was `true` while `{a: 1} <= {a: 1.0}` was `false`. All six now share one `eql?`-based relation: type-strict and self-consistent. The only observable change is that `==` no longer coerces `1` and `1.0` to equal, which was incidental `Hash#==` fallout, not a decision, and was already contradicted by the subset operators.

It takes its own release because it stands on its own — any two `nil`-containing `Namos` raised on `==` before it, grouping or not — and because it is the prerequisite that lets a `nil`-valued group round-trip through `==`, which `group_by` (0.20.0) relies on for its partition/`as_detail` idempotence.

0.20.0 (2026-06-15): group_by returns a Collection

`Namo#group_by(dimension)` splits a `Namo` into a `Namo::Collection`, partitioning the rows by the values of the given dimension. This completes the `Enumerable` coherence pass begun at 0.11.0 — it is the one `Enumerable` method that produces row-shaped output and was not included there.

Why this is a separate release, and why it’s here at all `group_by` was originally scheduled for 2.x. It comes forward to 0.20.0 *because* 0.18.0 built `Namo::Collection` — its return type. The dependency is the whole story:

The 0.11.0 `Enumerable` pass shipped every method that returns a `Namo`. `group_by` was excluded not as a deferral but because it was structurally blocked: it returns a *partition* — a keyed set of sub-`Namos` — and the right type for that didn’t exist. `{key => Array<Row>}` (Ruby’s default) isn’t in the `Namo` family, and `{key => Namos}` (a plain `Hash` of `Namos`) is a half-measure that doesn’t carry the aggregation surface a partition wants. The honest type for “a `Namo` split into named pieces” is `Namo::Collection`, which is also the type for “named pieces assembled into a whole.” Assembly and partition are the same structure reached from opposite directions; `Collection` is the structure, and `group_by` is its partition-side constructor.

So `group_by` could not land until `Collection` existed. 0.18.0 built `Collection`; 0.20.0 is the first release in which `group_by` is expressible. The two-release split (0.18.0 then 0.20.0) records the causality: `Collection`’s existence is what unblocks and brings forward `group_by`.

Behaviour `group_by(dimension)` returns a `Collection` whose members are the groups — one member per distinct value of `dimension`, each a `Namo` holding the rows that share that value. Because `dimension` is a pre-existing dimension of the data, it is retained in every member's rows (the split happens *along* the axis, it does not consume it). This is what makes the split round-trip exactly idempotent with `as_detail` on the same dimension:

```
namo.group_by(:symbol)                # => Namu::Collection, one member per symbol
namo.group_by(:symbol).summary(:close, reducer: :mean)  # mean close per symbol
namo.group_by(:symbol).as_detail(:symbol) == namo      # true – exact round-trip
```

Each group member is named by its group value (`find('BHP')` returns the BHP member), so the `Collection`'s `find` and the whole assembly API apply uniformly to a partitioned `Collection`. The group value is the member name — and because it was already a dimension in the rows, there is no extrinsic-key conflation: the name and the dimension value coincide by construction.

Grouping by a derived dimension. Because every operation in `Namo` treats data and derived dimensions equivalently, `group_by` does too: grouping by a formula is not a special case to be refused but a first-class use — `group_by(:value_score)` must work exactly as `group_by(:symbol)` does, or the unified-treatment principle would be false at precisely the operation where partitioning on a computed category (a score, an action, a class label) is most expected. When the named dimension is a formula, `group_by` materialises it first — through the same projection path that materialises any derived dimension (`self[*data_dimensions, dimension]`, 0.16.0) — so the grouped-by formula becomes a stored column and is dropped, while every other formula carries through live. The user neither knows nor needs to know which kind the dimension was; that indifference is the feature.

Materialise-and-drop is the only invertible choice. If a materialised dimension kept its formula alongside its stored values, the name would be both data and derived at once — the aliased state 0.16.0 exists to eliminate — and inversion would require a provenance ledger to recover which dimensions had been formulae. Dropping the formula keeps the namespace self-describing: a name's status is always recoverable from the object alone. The cost is that the grouped-by formula is no longer live in the members; a user wanting both the frozen grouping value and a live recomputation binds a second name before grouping, spending a column visibly at their own call site.

This collapses the round-trip into one uniform law over the whole dimension namespace: `namo.group_by(d).as_detail(d) == namo[*namo.data_dimensions, d]` for any `d`, data or derived. For a data dimension, `namo[*data_dimensions, d]` is just `namo` (projecting all columns is identity), so the exact-original round-trip is the degenerate instance of the general law against the materialised form. Parameterised formulae are the honest exception: they cannot be materialised without arguments, so `group_by(:sma)` raises the 0.17.0 `ArgumentError` at the materialisation site, the same way `values(:sma)` does — uniform across every materialisation point, not `group_by` treating them as second-class.

The `nil`-key case falls out cleanly from the use-site naming decision (0.18.0). If some rows have `dimension` missing or `nil`-valued, they form a group keyed by `nil`, and that member is named `nil`. There is no insertion guard to trip — `<<` accepts it — and the member is simply unfindable by `find` (you can't `find(nil)` meaningfully), but its rows still materialise into the detail view and round-trip correctly, because the grouping dimension (`nil`-valued for those rows, but present) is intrinsic. No rows are dropped, no sentinel is invented; the `nil` group is a first-class member that happens to have a `nil` name. This is exactly why 0.18.0 moved name-enforcement to use-site rather than guarding `<<`.

Relationship to the Enumerable pass With `group_by` landing here, the coherence statement from 0.11.0 completes: every `Enumerable` method that produces row-shaped output returns a `Namo`-family type. The subset methods (`select`, `reject`, `sort_by`, `uniq`, `partition`, `take/drop`) return `Namo` (or `[Namo, Namol]`); `group_by` returns `Namu::Collection`. The family is `Namo` and `Namu::Collection`; the rule is uniform across both.

Implementation


```

def group_by(dimension)
  collection = Collection.new
  source = derived_dimensions.include?(dimension) ? self[*data_dimensions, dimension] : self
  source.data.group_by{|row_data| row_data[dimension]}.each do |value, rows|
    collection << self.class.new(rows, formulae: source.formulae.dup, name: value)
  end
  collection
end

```

Each group is constructed as an instance of the receiver's class (subclass type preserved), carrying the surviving formulae, named by the group value so the Collection can key on it. The ternary is the materialise decision: for a derived dimension source becomes the projection `self[*data_dimensions, dimension]` (the grouped-by formula now a stored column and dropped, every other formula carried live); for a data dimension it is `self` untouched, allocating no projection. The members therefore carry `source.formulae`, not `@formulae` — after materialisation that is the *surviving* set, and carrying `@formulae` would re-introduce the dropped formula and break exclusivity. A parameterised formula raises inside the projection, at the 0.17.0 materialisation site, so `group_by` needs no enforcement of its own.

A block form (`group_by{|row| ...}` computing the group key from a block rather than naming a dimension) is a natural extension. Deferred consideration: the block form's groups are keyed by a computed value with no corresponding dimension in the rows, so the split round-trip would be the *assembly* case (the computed key must be injected as a dimension on `as_detail`), not the exact-idempotent case. Worth supporting, but the dimension-named form is the primary one and the one that round-trips cleanly.

Tests

- `group_by(:symbol)` returns a `Namo::Collection`.
- The Collection has one member per distinct value of `:symbol`.
- Each member holds exactly the rows matching its group value.
- Each member retains `:symbol` (the grouping dimension is not consumed).
- Each member is named by its group value (`find(value)` works).
- Each member carries the parent's formulae.
- Each member preserves the receiver's class (subclass type).
- `group_by` does not mutate the receiver.
- The uniform inversion law: exact on a data dimension (`group_by(:symbol).as_detail(:symbol) == namo`); against the materialised form on a formula dimension (`group_by(:value_score).as_detail(:value_score) == namo[*data_dimensions, :value_score]`).
- `group_by` on an empty `Namo` returns an empty Collection.
- Grouping by a formula: partitions by the formula's value; materialises the grouped-by formula into a data column and drops it from the members' derived dimensions; alters *only* the grouped-by formula, carrying a live dependent formula through unchanged; a parameterised formula raises `ArgumentError` at the materialisation site.
- `group_by(:sector)` where some rows have `nil :sector` produces a `nil`-named member holding those rows; no rows are dropped; the split round-trip still holds.
- `summary/detail` work on a `group_by`-derived Collection identically to an assembled one.

Documentation

- README section on `group_by` returning a Collection, with the round-trip example.
- Note that this completes the Enumerable coherence pass started at 0.11.0.
- Cross-reference to Collection (0.18.0) for the assembly side of the same structure.

0.21.0 (2026-06-25): Formulae extraction

The formula collection — the `{name => callable}` map every `Namo` carries — is promoted from a bare Hash to a `Namo::Formulae` object. This is behaviour-preserving: no new capability, and nothing a user observes

changes beyond the type of the formulae accessor itself. It exists to give the formula collection an object to accrete behaviour onto, exactly as Row was extracted empty-handed at 0.1.0 and later grew value semantics (0.10.0) and arity dispatch (0.15.0).

Why an empty extraction is its own release Formulae answers exactly the subset of Hash that `lib/namo.rb` and Row already send the formula collection — `[]`, `[]=`, `keys`, `key?`, `empty?`, `each`, `delete`, `merge`, `reject`, `dup` — so every operator, projection, and comparison behaves as before: each message it sent still works and returns an equivalent result. The proof is the prior test suite passing unchanged. Naming the collection now, before it carries behaviour, lets a later release add capability (module absorption, `[]-resolution`) onto a type that already exists, rather than threading a second extraction through the operators at the same time as the feature. Relocate first under a green suite; add capability second — the discipline of the Row (0.1.0) and Enumerable (0.11.1) extractions.

The normalisation seam The one behavioural change is in the constructor: incoming formulae are normalised to a Formulae. A Hash — what every internal construction passes, being the result of `Hash#merge/Hash#dup` in the operators — is wrapped; a Formulae — what the operators now pass, since `Formulae#merge/dup` return Formulae — passes through unchanged. No other call site changed; they all send messages Formulae answers. The public formulae accessor therefore returns a Formulae now, not a Hash, and the two tests that compared it to a literal `{}` assert emptiness instead — the formula set being empty is the behaviour they pin, the wrapper type is not.

Three deliberate choices

- `to_h` returns a copy (`@store.dup`), not the live store. The bare-Hash version could not leak its internals because there was no wrapper; the wrapper must not introduce the leak.
- `==/eq/?/hash` are names-only (`keys.sort`), matching Namo's own formula equality, where two sets with the same names but different proc *objects* are equivalent (`@formulae.keys.sort` in `===/eq/?/hash`). A full-store equality would be stricter than Namo means, and a landmine if `Namo#===` is later migrated onto `Formulae#===`.
- `merge/reject/dup` return a Formulae, and include Enumerable — with each yielding `[name, callable]` pairs — gives `map/select/any?` for free, matching Hash's Enumerable behaviour.

Out of scope Module absorption (`Formulae#<</>[]-resolution`); migrating the operators to pass Formulae throughout and dropping the constructor's Hash-acceptance branch; moving the left-wins merge rule into `Formulae#merge`; and migrating `Namo#===/eq/?` onto `Formulae#===`. Each is a capability or a tidy-up that the named type now makes possible; none belongs in the extraction itself.

Tests

- A Formulae unit test pinning the three deliberate choices and the Hash-compatible surface (`[]`, `[]=`, `keys`, `key?`, `delete`, `merge`, `reject`, `dup`, `each`, plus Enumerable `map/select`).
- Construction tests for the normalisation seam: a Hash given by keyword is wrapped in a Formulae; a Formulae given by keyword passes through unchanged (same object).
- The prior suite passes unchanged but for the two formulae-accessor assertions, now asserting emptiness rather than a literal `{}`.

Documentation

- No README change — the extraction is invisible at the API surface.

0.22.0 (2026-06-26): The derive resolver

Formula *invocation* moves out of `Row#[]` and into `Namo::Formulae`. `Formulae#[]` keeps its 0.21.0 meaning — it returns the stored callable — and a new `derive(name, row, namo, *arguments)` resolves a named

formula to its *value*. `Row#[]` stops calling the proc itself and delegates: it keeps only the fail-fast argument guard, the formula-vs-data dispatch (`@formulae.key?`), and the data fallback (`@row[name]`), handing the invocation mechanics to its collaborator. Behaviour-preserving at the usage level — `row[:revenue]`, `namo[:x]` = formula, values, and selections are all unchanged. Like the 0.21.0 extraction it enables the next release rather than changing this one: once `Formulae` owns resolution, 0.23.0's formulary methods are ordinary store entries that `derive` invokes with no new resolution path.

Why `derive`, not `call` The receiver is the collection, not the callable, so the verb names what the collection does — *derive a value via the named rule* — rather than borrowing `Proc#call`'s “invoke the receiver”. And the result is whatever the proc returns, not necessarily numeric, which a neutral verb leaves open. `Formulae#[]` returning the stored callable is storage access, not a `Namo` operation, so it carries no inconsistency with the dimension-value `[]` of `Namo` and `Row`: different domains, the same spelling.

What moved `derive` lifts the formula branch whole from `Row` — collection-scoped formulae are called (`row`, `namo`, `*arguments`), row-only formulae (`row`). The private `collection_scoped?` and `raise_unless_namo_context` move across with it (the latter now takes `namo` as a parameter rather than closing over `Row`'s `@namo`), and `required_parameter_count(name)` becomes a public property of the stored callable, since it is a fact about the formula, not about the row asking for it. `Row#[]` sources its argument-count check from `@formulae.required_parameter_count` and the fixed-vs-variadic distinction from `@formulae[name].arity`. `[]` and `[]=` are untouched — `[]` still returns the stored callable.

0.23.0 (2026-06-27): Formularies

A *formulary* is a reusable body of derived dimensions: a module, tagged by including `Namo::Formulary`, whose public instance methods are formulae. `namo.attach(OrderFlow)` makes those methods resolve as derived dimensions — through the same `Row#[]` → `Formulae` path, the same inspection vocabulary, and the same carry-through as `[]=` formulae. This is the module-based formula library the 2.x notes anticipated, brought forward as a real mechanism rather than a documented Ruby-include pattern, because `Namo::Formulae` (0.21.0) gave the formula collection an object to carry the behaviour and `Formulae#derive` (0.22.0) gave it the resolver. The dependency is the story: the type and the resolver existed, so the tier became an addition to `Formulae` rather than a new seam through the operators.

Attach copies into the store `attach` copies the formulary's `public_instance_methods(false)` into `@store`, each as a bound `Method`. `Formulae` holds **one** inert host (`@host ||= Object.new`) and extends each attached formulary into it, so every bound `Method` shares the same self-free receiver — `self` is a meaningless object, and a formulary method named like a `Namo` method can't shadow `Namo`'s own API (it's isolated on the host, not the `Namo`). From that point a formulary method is an ordinary store formula — indistinguishable from one assigned through `[]=`, and resolved by the unchanged 0.22.0 `derive`. So almost nothing in the tier is special-cased: `keys`, `key?`, `required_parameter_count`, `derive`, `dup/reject/merge`, and `Namo#derived_dimensions/==/requires_arguments?` all revert to their store-only forms and handle the formulary methods because they *are* store entries. `Namo::Formulary` (the marker) gates `attach` and `public_instance_methods(false)` selects the derivations — private helpers and inherited/`Object` methods are excluded with no per-method declaration, which is also why no `Namo` or `Object` API leaks as a dimension.

Two channels: `attach` and `class include` A formulary reaches a `Namo` two ways. `attach/<<` adds one to an individual `Namo` at runtime. Alternatively, `include` the formulary into a `Namo` subclass (`class TradingAnalysis < Namo; include OrderFlow; Namo#initialize` scans `self.class.ancestors.reverse`, and for each tagged module calls `attach(mod)` — so the class channel funnels through the same guarded `attach`, the same single host, the same store. The reverse is load-bearing: `ancestors` lists most-recent-include first, but `attach` is last-write-wins, so iterating in reverse makes the most-recently-included formulary win, matching Ruby's own method-resolution order. The contract: *the class include is honoured once, when each instance is constructed; everything after is attach*. Both copy, so `include-ing` into a class after an instance

exists doesn't reach that instance (re-build or attach), and a plain `Namo` — no subclass to include into — always takes `attach`. Routing the class channel through `attach` (not bare `@formulae.attach`) keeps it **guarded**: an included formulary colliding with an instance's data raises at construction, so exclusivity holds with no hole — the class channel is `attach`-like (an explicit “give this a formulary”), not constructor-like, so it earns the guard rather than the constructor's unguarded trust.

Why the store, not a live carrier A live design was built first and rejected: Formulae holding the modules in an inert carrier's ancestor chain and resolving them live, nothing copied. It was elegant — the chain is the list, most-recent-wins falls out of Ruby's `include` order — but it re-created the exact aliasing 0.16.0 spent a release eliminating. 0.16.0's exclusivity machinery drops a materialised formula on projection via `@formulae.reject`; a carrier can't be rejected one method at a time, so a projected formulary name materialised into a data column *and* lingered as a live method — `values(:sma)` reading the snapshot while `first[:sma]` recomputed off the projected-away inputs, the precise 0.16.0 bug. Patching it meant either data-first precedence in `Row#[]` (the alternative 0.16.0 explicitly rejected) or a propagating suppression set reimplementing the drop the store already does. Copying into the store sidesteps all of it: projection's existing `reject` drops a materialised formulary method for free, and the tier sits *inside* the exclusivity model instead of beside it. The cost is that `attach` is a snapshot, not a live link (below) — judged the smaller price, since value-liveness (recompute from each row's current data) is fully preserved and only “reflect a reopened module” is lost.

Inside the exclusivity model Because the methods land in the store, the 0.13.0/0.16.0 rules apply unchanged:

- **Most-recent / last-write wins.** A second `attach` of a colliding name overwrites the first; `[]=` and `attach` interleave as plain last-write-wins on a shared name — the same rule `[]=` already documents, now spanning both channels.
- **Data collisions raise on attach.** `Namo#attach` raises `ArgumentError` if a formulary method's name collides with a data column, rather than silently clearing it. The asymmetry with `[]=` (which silently clears) is justified by *intent and blast radius*, not by the `*/**` precedent (composition raises because two peer operands have no last-write order to appeal to — `attach`, like `[]=`, does have one, so that reasoning doesn't transfer). `[]=` names one column at the call site, so replacing it is the explicit intent; `attach` names a *module*, so a data collision is a name the user never typed, possibly across several columns at once, and silently destroying data is dangerous. So `attach` surfaces it: resolve explicitly first — contraction or rename — then `attach`. Formula-vs-formula collisions stay last-write-wins (that's re-attachment, not aliasing).
- **Materialise-and-drop on projection.** Naming a formulary method in a projection snapshots it into a data column and drops the formula through 0.16.0's `reject`, so all access paths agree and dimensions lists it once. This is the whole reason for the store choice.
- **Carry-through is free.** `dup/reject/merge` carry the formulae because they are store entries — no carrier to rebuild, no list to track.

Namo#=== via the marker (Path B) The 2.x note flagged that `===/eq/?/hash`, comparing formula names via `@formulae.keys.sort`, would wrongly equate a module-bearing `Namo` with a vanilla one once module formulae became real. The store approach answers it as Path B foresaw without changing those methods at all: formulary methods *are* store keys, so `@formulae.keys.sort` already counts them — a `Namo` with a formulary attached is not `===` to a vanilla one, while two `Namos` exposing the same names by different channels (`[]=` vs formulary) are `===`. The marker keeps it precise: only tagged modules are copied, so only intentional derivations become keys; the comparison adds names, never proc bodies — the 0.6.0 stance, extended for free.

attach, and << The verb is `attach` — already the ROADMAP's word for adding derivations to a `Namo`. `Namo#attach(M)` forwards to `Formulae#attach(M)`; both return `self` and alias `<<`, so `namo << OrderFlow` reads as attaching the formulary and chains. `<<` here is a plain alias for `attach`, scoped to formularies — not the earlier polymorphic / module-lacing operator, which is dropped, and not a row-append (`Namo` has

no single-row add; that stays with +). `Collection#<<` keeps its own meaning (member-add): it defines its own `<<`, so the inherited alias never reaches a `Collection` instance — `<<` adds a constituent appropriate to the receiver, a member to a `Collection` and a formulary to a plain `Namo`.

respond_to?(:call) in []= Folded in as a separable footgun-fix: `[]=`'s formula branch widens from when `Proc` to `respond_to?(:call)`, so a directly-assigned `Method` or `lambda` registers as a formula rather than broadcasting per-row. Arrays are unaffected — they don't respond to `call`. This is independent of the formulary tier (the tier doesn't go through `[]=`); it ships here as its own change.

Snapshot, not a live link `Attach` is a snapshot: it copies the module's public methods *as they are at attach time*. This is the one PORO affordance the store choice gives up, and it's worth stating precisely, because real Ruby is more dynamic than it might seem. A normal `include/extend` is a *live link* — methods resolve through the ancestor chain, so a method added to the module later, a redefinition, and a further module mixed into it afterwards are all picked up by objects that already include it. A copy tracks none of that: change the module after attaching, or include another module into it, and an already-attached `Namo` is unaffected — re-attach to refresh. A live-link mechanism (re-resolving against the module, or holding it by reference for the name set while snapshotting values under `freeze`) is possible later if wanted; 0.23.0 ships the snapshot. Bare-name resolution inside method bodies (buys for `row[:buys]`) and `freeze`-gated memoisation both remain 2.x — a formulary method indexes its `row` explicitly.

Tests

- **Formulae:** `attach` returns self, rejects an untagged module, and copies the formulary's public methods into the store (excluding a private helper); `<<` attaches as `attach` does; `key?`, `required_parameter_count`, and `derive` resolve the copied methods; most-recent-wins on re-attachment; a single shared host across several attaches (the bound `Methods` share a receiver); carry-through through `dup/reject/merge`.
- **Namo (runtime attach):** `derived_dimensions` lists the copied names (private helper absent, vanilla `Namo` empty); resolution through `values/Row/selection`; two-arity windowing and the missing-namo raise; a parameterised method omitted from bulk and raising bare, materialising through a row-scoped wrapper; re-attachment; an instance `[]=` formula shadowing a formulary method; `attach` raising on a formulary/data-name collision; carry-through through `selection`, `projection`, and `concatenation`; `projection` materialising-and-dropping a formulary name (listed once as data, snapshot read on per-Row access); `===` against a vanilla `Namo` and against the same names via `[]=`; `<<` attaching as `attach` does.
- **Namo (class include):** an included formulary's public methods listed as derived dimensions (private helper absent); resolution; most-recent-include wins on a name collision (MRO via `ancestors.reverse`); raising at construction on a formulary/data collision; an instance built before a later `include` not seeing it while a fresh instance and an `attach` do.
- `[]=` polymorphic dispatch widened: a `Method` and a `lambda` register as formulae; an array still broadcasts.

Documentation

- **README:** a “Formularies” subsection under **Formulae** — `attach` (and its `<<` operator form) and the marker, `private-hides-helpers`, the copy-into-the-store / snapshot model, the two channels (`attach` at runtime, `class include` at construction) and the build-time-vs-after contract, the raise on a data collision, materialise-and-drop, carry-through, and the forward-note that bare-name bodies and `freeze`-gated memoisation arrive later. The Polymorphic `[]=` section now says “a callable” (`proc`, `lambda`, or `Method`) rather than “a `proc`”.
- **COMPARISON:** a “Formularies” row — other tools reuse computation as external function libraries producing stored columns; `Namo`'s formularies are modules whose methods become first-class derived dimensions of the dataset.

0.24.0 (2026-06-29): Polymorphic <<

`Namo#<<` widens from the 0.23.0 formulary-only alias into an “append a constituent” operator. The verb is the same on both types — *append the constituent appropriate to the receiver* — but the constituents differ. For a base `Namo` they are formularies and rows: a `Module` attaches (0.23.0, unchanged), a `Hash` appends as a data row, a `Row` appends its underlying hash. For a `Collection` the constituent is a member: a `Namo` adds, a `Module` attaches (inherited), and a loose `Hash/Row` raises. `attach` stays the formulary-specific method, `Module`-only and unchanged; `<<` of a `Module` routes to it, so the collision-guard and formulary semantics carry over verbatim.

Three boundaries Each exclusion is load-bearing, not an omission. **No bare-callable arm:** formula *registration* is `[]=`’s job — it dispatches callable-vs-scalar and enforces data/formula exclusivity by eviction. `<<` is *append a constituent*, a different verb; a bare `proc` through `<<` would duplicate that path and reopen the eviction-vs-raise question. So `[]=` assigns a named thing, `<<` appends a constituent, cleanly split. **No whole-Namo arm:** `namo1 << namo0` does not mean “append `namo0`’s rows” — that is `+`. Excluding it keeps `<<`-as-append distinct from `+`-as-combine and, on a `Collection`, makes a `Namo` argument unambiguously a *member*. **A bare Hash is always a row:** formulae never arrive through `<<` as a bare `Hash` (they arrive as a `Module`, or via `[]=`), so a `Hash` at `<<` is unconditionally a data row — the one legislated disambiguation the dispatch rests on.

The deliberate divergence `Namo#<<` and `Collection#<<` diverge on exactly the `Hash/Row` arm: a base `Namo` appends the loose row, a `Collection` refuses it. A `Collection`’s `@data` is derived from its members and rebuilt on every member-add, so a loose row has no durable home — the next `<<` would erase it. The honest response is to refuse and redirect to member-add, rather than accept-then-silently-drop. This is recorded as intentional so it does not read as oversight.

A row-append on a base `Namo` is a data mutation: it does not pass through `[]=`, so it does not trigger formula-eviction. 0.24.0 shipped it *unguarded* on the reasoning that the data/formula exclusivity invariant was `[]=`’s to keep — but that left a reachable hole (see 0.24.1 below), so 0.24.1 guards it by raising on a name collision instead.

Collection formulae: no new mechanism A “collection formula” needs nothing built. It is an ordinary two-arity formulary method whose `namo` argument is a `Collection` and whose body reaches a collection view (summary, detail, members); it resolves through the existing derive path over the `Collection`’s detail rows. Member-awareness lives in the method body, not in any marker or second store. Both a `Namo::CollectionFormulary` marker and a member-axis resolver were considered and rejected: tagging is admission-control (keep non-formulae from becoming live dimensions), not level-classification, and a formula’s “level” is a per-method property of its body; while per-member-*output* computations are *summary*’s territory (member-in invocation), not derived dimensions at all. The convention is to name the argument `namo_collection`, which a plain `Namo` lacking members/summary self-enforces by `NoMethodError` — no guard needed.

Pivot / reshaping: a recorded decision Whether `<<` (or any operator) should grow a widening/pivot step was settled as *out*, and the reasoning is recorded so it isn’t revisited blind. The non-lossy two-axis aggregation is long-form summary generalised — in-grain, a separate future item. The widening-to-a-grid is *presentational*: it breaks dimension peerage (every key a dimension, none privileged), so it belongs at an output edge, not in the algebra. Neither is 0.24.0, and neither is `<<`.

Tests

- `Namo#<<`: a `Hash` appends as a data row (and returns the `Namo`); a `Row` appends its underlying hash; a whole `Namo` raises `TypeError` (that is `+`); a bare callable raises `TypeError` (that is `[]=`); a mixed chain interleaves rows and a formulary. The 0.23.0 `Module`-arm tests are unchanged.

- `Collection#<<`: a formulary attaches and its row-scoped method resolves over the detail rows; a collection-scoped method (the `namo_collection` convention) reaches members; a loose Hash and a loose Row each raise `ArgumentError` with the member-redirect message; a chain mixes member-adds with a formulary attach. The 0.23.0 member-add tests are unchanged.

Documentation

- README: the Formularies `<<` note widened — `<<` appends a constituent (Module, Hash, or Row) to a base `Namo`, chains, and raises on a bare callable or a whole `Namo`; the bare-Hash-is-always-a-row and row-append-does-not-evict notes. The Collection “`<<` and unnamed members” section gains the member/Module/raise dispatch and the deliberate Hash/Row divergence; a new “Collection formulae” subsection documents the `namo_collection` convention with no marker and no second store.

0.24.1 (2026-06-30): Guard row-append against a formula collision

0.24.0’s row-append arms (`<<` of a Hash or Row) went in unguarded, on the stated reasoning that the data/formula exclusivity invariant was `[]=`’s to keep. The reasoning had a hole. When an appended row carries a key matching an existing formula name *and* lands as `@data.first` — the mainline case of streaming the first row into an empty formulary-bearing `Namo` — the two access paths disagree: bulk `values` reads the raw datum (because `data_dimensions` keys off `@data.first.keys`) while per-Row access derives the formula (Row#[] checks `@formulae` first). That is precisely the split-brain the data-or-derived-never-both invariant exists to prevent.

So `add_row` now **raises** an `ArgumentError` naming the collision rather than admitting the row. The guard is symmetric with `attach`’s, and the asymmetry against `[]=` is deliberate: `[]=` names one column at the call site, so its scalar form *evicts* a same-named formula as your explicit intent; a row is a bulk append you didn’t name column-by-column, so refusing is the safe response, not silently evicting a formula on a data operation. Resolve it the same way as the `attach`-side collision — contract the formula away or rename — then append. Nothing else in `<<` changed.

0.25.0 (2026-06-29): **attach!** — forceful formulary attachment

`attach` (0.23.0) raises when a formulary method’s name collides with a data column, on the blast-radius reasoning: a formulary names a whole *module*, so a collision is a name you never typed — possibly several columns at once — and silently destroying data is dangerous. **attach!** is the explicit override. It **evicts** the colliding data columns first — deleting each exactly as assigning a proc through `[]=` clears a same-named data column — then attaches. The bang is consent: the blast-radius objection dissolves once the eviction is intentional, named, and bang-marked, so `attach` keeps the guard as the safe default and **attach!** is its deliberate forceful sibling — the same relationship `[]=` already holds between its silent eviction (one named column, explicit intent) and `attach`’s raise (a whole module, names you didn’t type). `<<` is unchanged: a `Module` still routes to the guarded `attach`, so the operator never evicts and **attach!** is reached only by explicit call — one safe operator, one explicit forceful method, mirroring `[]=` having no operator shortcut for its own eviction.

The driving case is refreshing a formulary over a *materialised snapshot*. Once a derived dimension has been projected into a stored column (0.16.0 *materialise-and-drop*), re-attaching the formulary that defined it collides with the now-data column; **attach!** re-attaches over it in one step, where `attach` would require an explicit contraction first.

One hazard has no counterpart in `[]=`’s scalar[]proc eviction; a formulary method that reads a column its own name evicts recurses on access. The eviction removes the datum, and exclusive storage leaves no shadowed value behind the formula, so `Row#[]` re-enters the formula rather than falling back to data. `[]=` rarely bites this way because the incoming proc is written at the call site, beside the name it replaces; a formulary is authored elsewhere, so a method reading its own evicted column is easy to attach without noticing. No recursion guard is added — `derive` resolves self-reference unguarded, as the formula mechanism has since

two-arity (0.15.0) and parameterised (0.17.0) formulae: Ruby doesn't guard method recursion either, and a cycle detector would tax every resolution to catch a bug the stack trace already names in a handful of frames.

0.26.0 (2026-07-04): **detach** — remove a formula or a formulary

`detach(constituent)` completes the formulary-side mutating vocabulary — `attach / attach! / detach` — and becomes the sanctioned formula-removal idiom, displacing the `namo.formulae.delete(:name)` reach-through in tests and docs. It sits on both `Formulae` (the worker) and `Namo` (which forwards and returns self), the same two-level shape as `attach`; `Collection` inherits it harmlessly. The dispatch is polymorphic, per the `[] / []= / <<` precedent: a `Symbol` deletes that formula from the store; a `Module` deletes each of its `public_instance_methods(false)` names; anything else raises `TypeError`. The `Formulae` worker owns the dispatch and the raises — `Namo#detach` needs no guard because there is no data-side interaction.

The settled decisions, several of them transplants of rules argued in earlier releases:

- **By-name, no provenance.** The store is a snapshot with no provenance ledger (0.23.0), so the `Module` arm means “delete the names this module contributes” — it deletes whatever *currently* holds each name, including a `[]=` formula that later overwrote a formulary method of the same name. This is honest by-name semantics, not a bug, and is pinned by a test.
- **The marker is required.** An untagged module raises `ArgumentError`, mirroring `attach` — you can only detach what could have been attached, and the check catches passing the wrong module at zero cost.
- **Absent name is a no-op returning self.** The `/` looseness rationale transplanted: detaching nothing leaves the original, so a cleanup step composes over any `Namo` without special-casing applicability. `Hash#delete` is silent on absent keys; same precedent.
- **Derived side only.** `detach` never touches `@data`. A `Symbol` naming a data dimension isn't in the store, so it's a no-op and the column survives; removing a data dimension is contraction's job (`namo[-:name]`). `detach` and contraction complement across two axes — data/derived and algebra/mutation: contraction is non-mutating algebra removing a data dimension with formulae carried, `detach` is mutation of the receiver removing a derived dimension.
- **No bang variant.** Bangs exist where a guard exists; removing a formula can't collide with data, so there is nothing for a `detach!` to force through — a principled asymmetry with `attach/attach!`, recorded the way the `no-add_row!` asymmetry was at 0.25.0.
- **No operator form.** `<<` appends constituents; there is no removal operator (`>>` would be cuteness over clarity), per the 0.14.0 governing principle — grow the operator surface only where composition can't express the operation. `detach` is a named administrative act, not algebra.
- **Single argument, chaining for multiples.** `attach` takes one module; `detach` takes one constituent. Multiple removals chain: `namo.detach(:a).detach(:b)`.
- **Host ancestry is untouched.** `extend` can't be undone, so a detached formulary's module remains in the inert host's singleton chain. Harmless and invisible: resolution is store-driven and the store entries are gone, so a re-attach restores resolution cleanly.

`Formulae#delete` is unchanged — it stays the store-level primitive `detach` sits on; only the reach-throughs in the tests migrate to `detach`. `Splat/multi-name detach` is out (chaining covers it; revisit only on demonstrated need).

Tests: `Formulae_test` gains a `#detach` describe (both arms return self; the `Symbol` arm removes the named formula and no-ops on an absent name; the `Module` arm removes exactly the formulary's method names, leaves a `[]=` name it doesn't contribute, and deletes a `[]=` overwrite occupying a formulary-method name — the no-provenance pin; `ArgumentError` on an untagged module; `TypeError` on a `String` and a `Hash`; chaining). `namo_test` gains a `#detach` describe (returns the `Namo`; removes a `[]=` formula from the queryable namespace; leaves others resolving; the `Module` arm empties `derived_dimensions`; a data dimension is a no-op; absent name no-ops; the two raises). `Collection_test` pins that a `Collection` inherits `detach`. The `#dimensions` “reflects mutation” test migrates `sales.formulae.delete(:revenue)` to `sales.detach(:revenue)`.

0.27.0 (2026-07-10): A block form on `summary` / `as_summary`

`Collection#summary(dimension, by:, reducer:)` reduces one dimension with one named reducer — `send(reducer)` over a single column. Per-member work that wants several columns at once, or a reduction that isn't a single-column `send` — a *pick-by-extremum* (“the close on the earliest-dated row”, an argmin that reduces no one column) is the case that forces it — is expressible with neither a single dimension nor a single reducer. An optional block supplies the per-member reduction the plain form can't name: it receives each member and returns that member's summary hash, and `summary` labels it and collects one row per member.

This is the warrant the 0.14.0 blocks on `*/**` carry — a block form earns its place exactly when the operation gives consideration to a dimension (here, a member) in isolation and the plain arguments can't express the consideration. It is one method growing a block, not a sibling: `summarise` was considered and refused — one name, one concept, the named reducer a shorthand for the block's general case rather than a second method beside it.

The pinned decisions:

- **dimension is optional (= nil).** `summary(by: :security){...}` is a block-form call with no column to name; without the default the missing positional would raise before the block was reached.
- **map, not flat_map.** `summary` reduces each member to exactly one row — one member in, one row out — where detail *expands* each member to all its rows. The opposite cardinalities are the two views' whole distinction, and the `map` / `flat_map` split is where it lives.
- **The member label wins.** `{by => member.name}` merges *last*, so a block returning a hash that carries the `by` key cannot rename the group — `by` is `summary`'s contract, not the block's.
- **Neither a dimension nor a block raises.** An `ArgumentError` names the omission rather than falling into the dimension branch and failing obscurely in `values(nil)`.
- **as_summary grows the block by delegation.** It forwards `&block` to `summary` and persists the result, so the capability lives in one place. The 0.18.0 persistence semantics are unchanged — a block-form view persists until the next `<<`, as the no-block form does.

The block is transient — never stored, never serialised — so the feature is serialisation-independent. It is the intermediate step toward a stored *group-scoped formula* (a per-group reduction as a named, reusable derived dimension, a third scope beside row- and collection-scoped), which is deferred: it needs a new formula scope, entangles with serialisation, and wants the 2.x freeze-gated caching to be affordable. Design it alongside those, not before.

Tests: `Collection_test` gains the block form returning one row per member merged with the member label; the member label winning over a `by` key the block returns; the neither-arg `ArgumentError`; the `map/flat_map` cardinality pin; and `as_summary` with a block setting the data and delegating. The no-block path is byte-identical, so the existing `summary/as_summary` tests are unchanged.

0.28.0 (2026-07-13): Positional-or-keyword label on `detail` / `as_detail`

`detail` took its origin label by keyword (`detail(by:)`), `as_detail` positionally (`as_detail(by = :member)`) — twin views of the same `Collection` with divergent calling conventions. The 0.18.0 rationale for the split, “`as_detail` carries no dimension, so its label is positional,” doesn't hold: `detail` carries no dimension either, so “no dimension” never distinguished the two. The real driver was only that `as_detail(:assembly)` reads well as the `group_by` inverse, `group_by(:symbol).as_detail(:symbol)`.

The resolution is the constructor's own precedent. `initialize(positional_data = nil, data: [], ...)` takes the same value positionally or by keyword, positional winning, rather than forcing a choice. Both `detail` and `as_detail` now take (`positional_by = nil, by: :member`) and resolve `by = positional_by || by`, so `as_detail(:assembly)` and `as_detail(by: :assembly)` are one call and the asymmetry dissolves.

Scope is `detail/as_detail` only — the methods whose label is the sole positional. `summary/as_summary` keep `by:` a keyword beside their dimension positional, mirroring the constructor exactly: the primary value (`data`) is the dual one, the modifiers (`formulae:`, `name:`) stay keyword. Here the primary value is the label.

It adds accepted call forms without changing any existing call's result — a minor, classed like the constructor widening that first introduced this convention: nothing was wrong before, downstream code can now depend on the keyword form, and the convention it conforms to was itself a minor.

0.29.0 (2026-07-15): Access-scoped materialisation

A formula that scans its parent asks for the same column once per row it visits, so materialising one derived column re-materialises its inputs n times over — an $O(n^2)$ that is pure redundancy within a single access. At a few hundred rows this is the unmeasured cost the 0.7.0 note anticipated; at the hundreds of thousands an end-of-day equities set reaches it is $\sim 10^{11}$ operations and unusable. `values` now opens a per-access scope that a nested `values` call reuses rather than reopening, `values_for` memoises its column build into `@materialised[dim] ||= ...`, and the scope is discarded when the top-level access returns. Each column materialises once per access and is reused; the parent-scanning $O(n^2)$ collapses to $O(n)$ — measured against a `return→sma→signal` chain, Row allocation drops from n^2+n to $2n$.

This refines the 1.x discipline rather than breaking it. The pure-live rule was stated as “recompute from current state on every access, no memoisation anywhere” — a proxy for its real content, which is *no staleness*. The precise invariant is **no derived value outlives the access that produced it**. A per-access scope caches only within the resolution of one access, where the state cannot change, and discards the cache before control returns to the caller — so every access still recomputes from current state and nothing can go stale. It is the evaluator hoisting a loop-invariant, the value the formula's author would otherwise bind to a local, not caching in the sense the freeze work concerns. The `attr_accessor :data` back door that defeats a generation counter is irrelevant: a value that never outlives its access never needs to observe mutation. Access-scoped materialisation is the freeze-free floor beneath the freeze-gated 2.x ceiling — the same `values_for` materialisation seam at a shorter lifetime, one access here versus until-unfreeze there — so the two compose, and neither runs mutation-observing invalidation.

Scope is the `values` path only — `values`, `coordinates`, `to_h`. Iteration through `Namo::Enumerable` (each, select, map, group_by) builds Rows straight off `@data` and does not open the scope, so a parent-scanning formula consumed by iteration is not yet covered. Bringing it under the scope is the same one-line wrap at each traversal entry, deferred here pending an iteration-heavy real workload: a `return→sma→signal` microbenchmark confirms the `values`-path scope restores $O(n)$, but the `Enumerable` churn stays out until a workload that consumes derived columns through iteration shows it insufficient. The residual intra-access cost it would not touch — a formula genuinely rescanning a window per row — is not redundancy, and no scope removes it; that needs sorted/indexed windowed access, a separate primitive.

It adds no observable behaviour — same results, same liveness, faster — so it is a minor, and contract-neutral: the scope materialises whole columns however a formula is expressed, opaque procs exactly as they are.

0.30.0 (2026-07-15): Single-pass group_by assembly

`group_by` assembled its `Collection` one member at a time, and `Collection#<<` rebuilds `@data = detail.data` over all members-so-far on every append (0.18.0's rebuild-on-`<<` model). Partitioning into g groups therefore re-materialised the detail g times — $O(\text{rows} \cdot \text{groups})$, the same quadratic shape as 0.29.0's parent-scanning formula but in the assembly path rather than the `values` path. At the group counts an end-of-day equities set reaches — thousands of securities — it is the dominant cost. `group_by` now maps the groups to members and hands them to a single `<<`, so the detail materialises once whatever the number of groups; the assembly is $O(\text{rows})$.

It stays inside 0.18.0's resolved model rather than reopening it: `<<` still rebuilds `@data` eagerly, the snapshot is still taken at append, and an `as_summary/as_detail` view still persists until the next `<<`. What changes is only *how many* appends `group_by` makes — one, not g . It adds no observable behaviour and is contract-neutral, a minor. The case 0.18.0's “`<<` is the only mutator” premise does not cover — mutating a member *after* it is added, which the eager `@data` would not reflect — is untouched here; reconciling it with the per-access intent (a truly lazy `@data`, or a `member→Collection` propagation) remains the live-composition question, deferred.

Summary

The set operators (+, -, &, |, ^), the comparison operators (==, ===, eql?, <, <=, >, >=), and the composition operators (*, **, /) — * and ** taking optional blocks for custom match refinement — together with selection (exact, array, range, proc, regex), projection, contraction, formulae (one-arity row-scoped, two-arity collection-scoped, and parameterised receiving arguments at access time, mixing freely), polymorphic assignment via []= (a callable — proc, lambda, or Method — registers a formula, scalar broadcasts to every row, exclusive storage either way), data/formula exclusivity carried through projection (naming a derived dimension materialises it and drops the formula; omitting it carries the formula live) and composition (* and ** refuse a data/formula name collision), the full inspection vocabulary (dimensions, data_dimensions, derived_dimensions, coordinates, values, to_h), Row value semantics (==, eql?, hash), the subset-returning Enumerable methods (select, reject, sort_by, first, last, take, drop, take_while, drop_while, uniq, partition) returning Namos, a constructor that takes data positionally or by keyword and carries an optional name:, `Namo::Collection` — a hierarchical aggregate of named member Namos, assembled with << and queried through summary/detail views with lazy detail materialisation, `summary/as_summary` taking an optional block for a per-member reduction beyond a single named reducer — and `group_by`, the partition-side constructor that splits a `Namo` into a `Collection` (the mirror of assembling one), formularies (`Namo::Formulary`) — reusable modules of derived dimensions attached to a `Namo` at runtime through `attach/<<` (or `attach!`, the forceful sibling that evicts a colliding data column rather than raising) or mixed into a subclass through `include`, and removed by `detach` — the mutating family is `attach/attach!/detach` — resolving as first-class derived dimensions — and the polymorphic << operator that appends the constituent appropriate to its receiver (a formulary or a data row to a base `Namo`, guarding a row against a data/formula name collision; a member to a `Collection`), give `Namo` a complete vocabulary for working with a single dataset, combining datasets that share the same dimensions, combining or decomposing datasets with different dimensions, composing named datasets into a queryable whole, partitioning one back into named pieces, and drawing on reusable libraries of derived dimensions, with Rows that behave correctly as Ruby values, cross-row computation that reflects the live state of the `Namo` it's asked through, and analytical chains that stay closed through filtering and ordering. The next phase is the 1.0.0 stable release.

1.0.0: Stable release

The 1.0 release includes everything through 0.27.0:

- Selection (exact, array, range, proc, regex), projection, contraction.
- Single-row formulae, two-arity formulae, parameterised formulae.
- The set algebra (+, -, &, |, ^).
- The composition algebra (*, **, /).
- The comparison algebra (==, ===, eql?, <, <=, >, >=), at both `Namo` and `Row` levels.
- Block forms on every operator that pairs rows.
- The inspection vocabulary (dimensions, data_dimensions, derived_dimensions, coordinates, values, to_h).
- Enumerable methods returning Namos for subset operations.
- Constructor widening (keyword `data:` and `name:`) and polymorphic []=.
- `Namo::Collection` for hierarchical aggregates, reachable by assembly (<<) or partition (`group_by`).
- `summary/as_summary` taking an optional block for a per-member reduction beyond a single named reducer.
- `group_by` returning a `Collection`, completing the Enumerable coherence pass.
- `Namo::Formulary` as the formula collection's own type, extracted from a bare Hash, with `derive` as its value resolver.
- Formularies — reusable modules of derived dimensions a `Namo` attaches, resolved live and carried through the algebra.
- Polymorphic << — appends a constituent: a formulary (`Module`), a row (`Hash` or `Row`) on a base `Namo`; a member (`Namo`) or formulary on a `Collection`.
- `attach!` — the forceful sibling of `attach`, evicting a colliding data column rather than raising when a formulary is attached over it.

- `detach` — removes a formula by name or a formulary's methods by module, completing the `attach/attach!/detach` family.

This is the correct, tested, conservative foundation. No metaprogramming magic, no `method_missing`, no `instance_eval`. Formulae work via `e[:name] = proc{|row| row[:close] / row[:book_value]}` — clear, explicit, proven.

1.0 is the set of features that are well-understood, thoroughly tested, and unlikely to change.

Estimated performance: ~0.3s for a 2,000 row daily trading screen with indicators and scoring. Pure Ruby, no native dependencies. Adequate for interactive use and daily batch jobs. Not suitable for datasets over ~50,000 rows without patience.

1.1: Benchmarking suite

Performance and stress-testing infrastructure. Establishes baselines for every feature so that subsequent optimisations can be measured against concrete numbers.

The suite should cover:

- Construction: `Namo.new` from N rows, measuring hash ingestion and dimension/coordinate inference.
- Selection: single value, array, range, proc, regex, at various dataset sizes.
- Projection and contraction: `[]` dispatch overhead.
- Formulae: row-scoped resolution, formula chains (A references B references C), two-arity formulae, parameterised formulae.
- Enumerable: `each`, `map`, `select`, `reduce`, `max_by` — measuring Row object creation per iteration.
- Set operators: `+`, `-`, `&`, `|`, `^` at various dataset sizes.
- Composition: `*` across different dimension overlaps.
- Collection: `<<`, `find`, `summary`, `detail` at various member counts and per-member row counts.

Each benchmark at multiple scales — 100, 1,000, 10,000, 100,000, 1,000,000 rows — to show where curves bend and where Ruby's overhead dominates.

Use `benchmark-ips` for iterations-per-second measurements. Results should be recorded and versioned so regressions are visible across releases.

The 1.1 numbers become the baseline for 2.x comparisons.

1.2: Loaders

Optional requires for common data sources. No loader loaded by default. No dependencies added to `Namo` core.

```
require 'namo/loaders/csv'
require 'namo/loaders/sequel'
require 'namo/loaders/json'
```

Each extends `Namo` with a class method (`from_csv`, `from_sequel`, `from_json`). Loaders are sugar — the constructor already takes arrays of hashes, so loading is always possible without them.

2.x: Bare names and Ruby-side optimisation

Theme: the expressive leap.

The 1.x / 2.x boundary is a discipline boundary as much as a feature one. **1.x is correctness and coherence, and pure-live throughout** — every derived view (`values`, `coordinates`, `to_h`, and `Collection`'s `find/summary/detail/lazy @data`) recomputes from current state on every access, and no derived value outlives the access that produced it. Liveness is unconditional; there is no staleness to reason about. The one thing 1.x does memoise (0.29.0) lives strictly inside a single access — a loop-invariant hoist that is discarded

before the access returns, so it cannot span a mutation and needs no invalidation. **2.x is the performance phase**, where *cross-access* memoisation lands across the board, alongside bare names, `method_missing`, `DefineAccessors`, and `Finite`. Cross-access caching is introduced uniformly here rather than piecemeal in 1.x, so the rule is simple to state and simple to trust.

Estimated performance: `method_missing` on Row adds ~5-10% overhead versus 1.x on first access per dimension per Row. The `DefineAccessors` optimisation (below) recovers most of this by defining real methods on first access. After warm-up, 2.x should be within ~2-3% of 1.x. Pure Ruby performance tuning against the 1.1 benchmarking suite may yield further gains. All estimates — actual numbers depend on the 1.1 baseline.

From 2.0, require 'namo' loads bare name resolution by default. The full experience is the default. Users who want the 1.x explicit style can opt out:

```
require 'namo'           # 2.x default: bare names included
require 'namo/core'      # opt out: 1.x explicit style, no method_missing on Row
```

Bare name resolution

Row resolves dimension names and formulae as bare method calls via `method_missing`, delegating to `self[name, *args]`. This eliminates hash access syntax from formulae:

```
# 1.x style
e[:price_to_book] = proc{|row| row[:close] / row[:book_value]}
```

```
# 2.x with bare names via `method_missing` on Row
e[:price_to_book] = proc{|row| row.close / row.book_value}
```

```
# 2.x with bare names via `instance_eval` for arity-0 procs
e[:price_to_book] = proc{ close / book_value }
```

Bare names also work in cross-row formulae. Fixed dimension references resolve against the current Row. Variable field names (passed as parameters) use `Symbol#to_proc` via `&field`:

```
e[:sma] = proc do |row, namo, field, period|
  window = namo[symbol: symbol, date: ..date].last(period)
  window.sum(&field) / window.count.to_f
end
```

Here `symbol` and `date` are bare names resolving against the Row. `&field` uses `Symbol#to_proc` to send the field name to each Row in the window.

Resolution order in `Row#method_missing`: formulae first, then row data, then super (raises `NoMethodError` for typos).

`method_missing` on Namo for formula assignment

Namo resolves formula assignment as bare method calls via `method_missing`, catching `name=` calls:

```
e.price_to_book = proc{ close / book_value }
e.sma = proc do |row, namo, field, period|
  # ...
end
```

Module-based formula libraries (documented pattern)

Partly shipped at 0.23.0. The *formulary* tier shipped the explicit form of this: a module tagged `Namo::Formulary`, attached with `namo.attach(M)`, whose methods take `row` (and `namo`) explicitly. What remains 2.x is the *bare-name* body shown below — methods that reference `close/symbol/date` directly rather than indexing `row` — and the `include-into-subclass`

ergonomics that bare-name resolution via `method_missing` makes possible. 0.23.0 is the registration channel; 2.x is the bare-name reading on top of it.

With bare name resolution in place, formulae can be defined as plain Ruby methods in modules and included into Namo subclasses. This is not a Namo feature — it's standard Ruby. It's documented here as a recommended pattern:

```
module Indicators
  def sma(row, namo, field, period)
    window = namo[symbol: symbol, date: ..date].last(period)
    window.sum(&field) / window.count.to_f
  end

  def change
    close - open
  end

  def earnings_yield
    eps / close
  end
end

module Scoring
  def value_score
    return 0 unless pe && pe > 0
    pe < 10 ? 2 : pe < 15 ? 1 : pe < 25 ? 0 : -1
  end

  def total_score
    value_score + book_score + yield_score + momentum_score + trend_score
  end

  def action
    s = total_score
    s >= 5 ? 'BUY' : s >= 2 ? 'WATCH' : s >= 0 ? 'HOLD' : total_score <= -2 ? 'SELL' : 'AVOID'
  end
end

class TradingAnalysis < Namo
  include Indicators
  include Scoring
end
```

Scoring strategies are swappable via dependency injection.

Namo#=== revision under module-based formulae

Resolved at 0.23.0 (Path B). With the formulary tier real, `===/eq?/hash` now compare `derived_dimensions.sort` — the union of `@formulae.keys` and the attached formularies' public method names — so a formulary-bearing Namo is no longer `===` to a vanilla one, while `[]=-vs-`formulary parity holds. The marker (`Namo::Formulary`) is what makes the union precise. The analysis below is retained as the rationale; it now describes shipped behaviour for attached formularies, with the residual gap being only the bare-name `include-into-subclass` channel once that arrives.

When module-based formulae become a real pattern (rather than a documented one), `Namo#===` (and

Namo#eq1?, and Nam0#hash) must be revised to enumerate module-defined method names alongside @formulae.keys. The 0.6.0 implementation compares formula *names* via @formulae.keys.sort — which is correct for []=registered formulae, but module-defined formulae don't enter @formulae. They're plain Ruby methods on the included module, discovered by Row's method_missing through the normal method-lookup chain. So a TradingAnalysis instance with include Indicators would have @formulae.keys == [] even though it has the full Indicators analytical surface; under 0.6.0's ==, two such instances are correctly == to each other (same dimensions, same — empty — @formulae.keys), but they'd also be == to a vanilla Nam0 of the same dimensions, which is wrong: a vanilla Nam0 has no indicators.

The starting point

```
# 0.6.0 implementation
def ==(other)
  return false unless other.is_a?(Nam0)
  dimensions.sort == other.dimensions.sort &&
    @formulae.keys.sort == other.formulae.keys.sort
end
```

Correct for the case where all analytical structure flows through @formulae. The gap is everything that lives on the class instead of the instance.

Alternatives considered Three candidates were on the table:

Path A: Check class identity. == also requires self.class == other.class, matching eq1?'s class-strictness. Same class implies same included modules implies same module-defined formulae, by transitivity. Simple to implement but rejects useful equivalences — two Nam0s from different classes with the same queryable surface (one via []=, one via include) would not be == even though they behave identically as analytical artefacts.

Path B: Duck-type the queryable surface. == compares the full set of names queryable on each Nam0, regardless of where each name's definition lives. Storage dimensions, instance @formulae keys, and class-level analytical method names all contribute. Two Nam0s with the same queryable surface are == regardless of class or how their formulae were registered. Matches the design philosophy's unified-treatment principle: care about what's queryable, not where the definition lives.

Path C: Class as proxy for analytical surface. Same code as Path A, framed differently — class identity stands in for the union of all analytical methods. This is just Path A with different rationale; the behaviour is identical and the same objection applies.

Path B is the right answer. It honours duck-typing across class boundaries, handles the []=vs-include mixture correctly, and doesn't change behaviour when users refactor a formula from per-instance assignment to a module method.

Path B implementation sketch

```
def ==(other)
  return false unless other.is_a?(Nam0)
  queryable_names == other.send(:queryable_names)
end

private

def queryable_names
  (dimensions + @formulae.keys + class_formula_names).to_set
end
```

`to_set` gives order-insensitive comparison; `.sort` would also work and avoids the `Set` allocation. Either is fine — pick when implementing. The equivalent changes to `eq?` and `hash` follow the same shape: substitute `queryable_names` for `@formulae.keys.sort` everywhere they appear.

What `class_formula_names` returns The open design question. Two candidate mechanisms:

Implicit identification. All public instance methods on the `Namo` subclass that aren't inherited from `Namo` itself are treated as formulae. Simple, no declaration overhead. Cost: conflates analytical methods with any custom method — a `def to_s` override on the subclass would be misread as a formula.

Explicit declaration. Modules opt into being treated as formula libraries, either by a marker module (include `Formula::Module`) or per-method declaration (`formula :sma`). Cost: declaration overhead. Benefit: precise — only intentionally-analytical methods count.

The implicit approach is more ergonomic; the explicit approach is more robust. Choose based on what the module-based formula mechanism actually looks like once it's implemented.

Decided at 0.23.0: explicit, via a marker module. A formulary is recognised by `include Namu::Formulary`, and its derivations are exactly its `public_instance_methods(false)` — so a `private` helper or an inherited/`Object` method is never misread as a formula, the precision the explicit approach buys. The union of names is read off the carrier's tagged modules; `class_formula_names` for the bare-name `include-into-subclass` channel will read the same marker off the subclass's ancestors when that channel arrives.

Body equality is not part of any equality operator 0.6.0 already settled this for the instance-level case: `==`, `eq?`, and `hash` all compare formula *names* (`@formulae.keys.sort`), not the procs themselves. The Path B revision extends that stance to module-defined formulae — adding more names to compare, not adding proc-body comparison. The rationale is the same as at 0.6.0 design time:

1. **Proc identity isn't function identity.** Two interactively-built `Namos` defining `:doubled = proc{|r| r[:x] * 2}` separately have different proc objects. Comparing by proc identity would call them unequal despite identical behaviour. Users would be surprised.
2. **`==` is a pattern-match operator, not a behaviour-identity operator.** Pattern-matching asks “does this candidate fit this pattern?” The pattern is the analytical shape — what names are queryable. `eq?` makes the same call for the same reason: in Ruby there's no cheap, reliable way to compare two procs for behavioural equivalence, so neither operator pretends to.

The deeper question of true behavioural equivalence (AST comparison, `iseq` comparison, DSL normal forms, etc.) is its own significant project; see the “Proc equivalence options” note in the project memory for the landscape. Until one of those paths lands, key-based comparison is the honest pragmatic position.

When this revision lands The attached-formulary case landed at 0.23.0, when `attach` made module-defined formulae a real registration channel. The residual case — bare-name methods included into a `Namo` subclass and discovered through `method_missing` — lands with bare-name resolution in 2.x; `class_formula_names` reads the same marker off the subclass's ancestors. Until then, the 0.23.0 union over store keys and attached formularies handles every shipped channel.

DefineAccessors optimisation

Inspired by Hashie's pattern — lazily define real methods on first access to avoid repeated `method_missing` overhead:

```
class Row
  def method_missing(name, *args)
    if @formulae.key?(name)
      define_singleton_method(name) do |*a|
        resolve(name, *a)
      end
    end
  end
end
```



```

        end
        send(name, *args)
      elsif @row.key?(name)
        define_singleton_method(name) { @row[name] }
        @row[name]
      else
        super
      end
    end
  end
end
end

```

First access fires `method_missing` and defines a real method. Every subsequent access is a direct method call. Consider when profiling shows `method_missing` as a bottleneck.

Hashie as prior art, not a dependency

Hashie (particularly `Hashie::Mash`) is the primary prior art for method-style hash access and coercion in Ruby. Namo's Row and the coercion module design draw from Hashie's patterns. However, Hashie should not be used as a dependency or as internal storage for Namo:

- Namo already handles method-style access through Row's `method_missing`. Wrapping internal hashes in `Hashie::Mash` would add a second `method_missing` layer — double the dispatch overhead for no additional capability.
- Hashie does key normalisation (string/symbol indifference) and deep conversion of nested hashes. Namo doesn't need either — keys are always symbols and data should be flat.
- Hashie has a documented history of subtle breakage, particularly around `respond_to_missing?` interactions with other gems and method name collisions with Ruby built-ins.
- The performance cost of Mash allocation over plain Hash allocation is measurable.

Pure Ruby performance tuning

Profile against the 1.1 benchmarking suite. Identify and address hot paths within pure Ruby: Row object allocation, formula chain resolution, selection dispatch. No native code — just better Ruby.

Caching / memoisation (opt-in via freeze, transparent)

1.x is pure-live: values, coordinates, `to_h`, and Collection's `find/summary/detail/lazy @data` all recompute on every access. 2.x adds memoisation as a performance layer, with two non-negotiable properties:

1. **Transparent.** Memoisation never changes observable behaviour, output, or liveness. Same inputs produce same outputs with the same currentness guarantees; only the cost differs. A user never has to reason about whether a cached value is stale — that is the cache's problem, and if it can't be guaranteed correct, the cache does not engage.
2. **Correct by construction, opt-in via freeze.** Caching engages only on **frozen** instances. An unfrozen Namo or Collection can mutate (`[]=`, `<<`, `@data/@members` changes), so its derived views could go stale under a cache — therefore an unfrozen instance always recomputes live, exactly as in 1.x. A frozen instance cannot change, so a cached derived value stays correct forever. The rule:
 - **Unfrozen** → pure-live, always recompute, no caching. (The 1.x default, unchanged.)
 - **Frozen** → caching may engage, because immutability guarantees freshness.

This makes memoisation genuinely optional without a flag, a mode, or a config the user must learn. The user opts in by doing the thing they would do anyway when finished mutating — `freeze` for sharing, hash-keying, or thread-safety (see the Mutability open question) — and the caching rides along for free. A user who never freezes never caches and never thinks about it. The optionality is about *whether caching engages* (the freeze

decision), never about *whether a cache is correct* — a cache only ever exists on immutable state, so stale-cache wrong answers are impossible by construction.

Scope of the 2.x caching work:

- **Namo inspection aspects** — memoise values/coordinates/to_h on frozen Namos (the “wholesale aspect caching” the 0.7.0 note defers to here).
- **Collection views** — memoise find/summary/detail/@data materialisation on frozen Collections. This is the memo-and-invalidation apparatus that was deliberately *removed* from the 0.18.0 implementation; it returns here, where it is measured against the 1.1 benchmark suite and justified — and where freeze makes invalidation-on-mutation unnecessary (a frozen Collection never mutates, so there is nothing to invalidate).

Contingency: this work is gated on the **freeze semantics**, currently in the Open Questions section rather than a scheduled release. Caching attaches to freeze, so resolving the freeze semantics is a prerequisite for the caching workstream. If freeze lands before 2.x’s performance phase, caching can attach to it directly; if the freeze semantics are still open when this work begins, settling them comes first.

Bare-name ergonomics on Collections

`group_by(:symbol)` returns a `Namo::Collection` from 0.20.0; the return type is settled in 1.x. What 2.x adds is the bare-name resolution that makes aggregation over a Collection’s members read cleanly:

```
# 1.x – works, but explicit
namo.group_by(:symbol).summary(:close, reducer: :mean)
# => Namu with {symbol:, close:} rows

# 1.x – explicit member-wise computation
namo.group_by(:symbol).members.map{|n| n.values(:close).sum / n.count}

# 2.x – bare names make member-wise computation read cleanly
namo.group_by(:symbol).members.map{|n| n.close.sum / n.count}
```

The return-type change (`group_by` → `Collection`) is *not* a 2.x concern — it landed at 0.20.0, gated on `Collection` at 0.18.0. 2.x only contributes the bare-name reading on the resulting members. The conceptual model was already unified in 1.x: assembly and partition both produce a `Collection`; bare names make the post-partition computation as ergonomic as the pre-partition selection.

Finite module

A module that includes `Enumerable` and adds `last` and `reverse_each` for finite collections. Default implementation uses `entries`; `Namo` overrides for performance by going straight to `@data`.

```
module Finite
  include Enumerable

  def last(n = nil)
    n ? entries.last(n) : entries.last
  end

  def reverse_each(&block)
    return enum_for(:reverse_each) unless block_given?
    entries.reverse_each(&block)
  end
end
```

Finite is a standalone concept — potentially a separate gem. Any finite `Enumerable` can include it.

Namo includes `Finite` and overrides `last` to go straight to `@data`, wrapping results in Rows. The Row constructor's `@namo` parameter (from 0.15.0) is passed through so `Finite`-wrapped Rows participate in two-arity formulae just like rows from each:

```
def last(n = nil)
  if n
    @data.last(n).map{|row| Row.new(row, @formulae, self)}
  else
    Row.new(@data.last, @formulae, self)
  end
end
```

`Finite` lands in 2.x rather than 1.x because it's primarily a performance optimisation — the 1.0 `last(n)` works correctly via the `Enumerable` default, just less efficiently. 2.x is the dedicated performance phase, and `Finite` belongs there.

Finite as a separate gem

Extract the `Finite` module into a standalone gem for use by any finite `Enumerable`, independent of `Namo`. Lands after the in-`Namo` version is stable.

3.x: DSL, columnar storage, and C acceleration

Theme: the optimised engine.

Estimated performance: columnar storage accelerates selection-heavy workloads — scanning one contiguous array versus touching every row's hash. C extension via `RubyInline` targets the remaining hot paths (iteration, row creation). Combined, the 2,000 row daily screen should drop from ~0.3s to ~0.1s. Selection on large datasets (100,000+ rows) should see the biggest improvement from columnar storage. The C extension adds ~4x speedup on tight numeric loops but the bottleneck is object allocation, not arithmetic — expect modest overall gains from C alone. All estimates.

From 3.0, require `'namo'` loads both bare names and the DSL block syntax by default. Users can opt out at any granularity:

```
require 'namo'           # 3.x default: bare names + DSL
require 'namo/core'      # 1.x explicit style only
require 'namo/bare_names' # bare names without DSL
require 'namo/dsl'       # DSL block syntax without bare names
require 'namo/dsl/bare_names' # both (same as default, explicit about it)
```

The DSL and bare names are independent features on different objects — the DSL uses `method_missing` on a builder to catch formula registrations, while bare names use `method_missing` on `Row` to resolve dimension access. The DSL block works without bare names (procs use explicit `row[:close]` inside), and bare names work without the DSL block (formulae registered via `e.name = proc{}`). In practice, most users will want both.

DSL block syntax

A minimal-ceremony specification syntax using a builder with `method_missing`:

```
namo do
  sma           ->(field, period) { window(field, period).mean }
  change       -> { close - open }
  price_to_book -> { close / book_value }
  earnings_yield -> { eps / close }
  golden_cross  -> { sma(:close, 20) > sma(:close, 50) }
  value_score   -> { pe < 10 ? 2 : pe < 15 ? 1 : pe < 25 ? 0 : -1 }
```

```

    total_score    -> { value_score + book_score + yield_score + momentum_score + trend_score }
    action         -> { total_score >= 5 ? 'BUY' : total_score >= 2 ? 'WATCH' : total_score >= 0 ? 'HOLD' : total
2 ? 'SELL' : 'AVOID' }
end

```

The builder's `method_missing` catches each name-lambda pair. No `def`, no `end`, no `module`, no `class`, no `e.name = proc`. Just names and computations.

Columnar storage

A columnar backend (hash of arrays) to accelerate selection and aggregation for pure-Ruby workloads. The design uses a mezzanine pattern — `Namo` delegates to `Namo::RowStore` or `Namo::ColumnStore` with the same public API.

This section is about *internal storage*, not about adding new public methods. `to_h` and `values` (both producing columnar output) already exist from 0.7.0; what 3.x changes is the cost of computing them. With row-oriented storage (current), `to_h` builds the columnar hash on demand by iterating rows. With columnar storage, `to_h` returns the internal representation directly. Same output, different cost.

The columnar form is exactly the shape that `to_h/values` produces — `{symbol: ['BHP', ...], close: [42.5, ...]}`, a hash of arrays. Columnar storage means accepting that shape as input and using it as the internal representation. The round-trip is clean: `Namo.new(namo.to_h)` reconstructs from columnar output.

Development order:

1. Default (current): always `RowStore`, no choice, no configuration.
2. Choose at instantiation: `Namo` detects the input shape (array of hashes vs hash of arrays) or accepts an explicit `storage` keyword.
3. Dynamic: one primary layout with the other cached on demand, cache invalidated on mutation.

```

# Row-oriented (array of hashes, current)
namo = Namo.new([{:symbol: 'BHP', close: 42.5}, ...])

```

```

# Column-oriented (hash of arrays) – the shape to_h produces
namo = Namo.new({symbol: ['BHP', ...], close: [42.5, ...]})

```

When DuckDB is the backend (4.x), neither layout matters — the data lives in DuckDB. Columnar storage only matters when `Namo` is doing the computation itself in pure Ruby.

C extension via RubyInline

Optional C acceleration for hot paths via `RubyInline` (`zenspider`). Targets: selection (linear scan over contiguous array), iteration (row creation), and storage layout. Formula resolution stays in Ruby.

Pure Ruby is the default. C extension is an optional require. The API doesn't change:

```

require 'namo'
require 'namo/ext' # optional C acceleration, falls back to pure Ruby if unavailable

```

Coercion modules

Two forms of coercion, both declared as includable modules:

Ingestion coercion — fix types from CSV/JSON:

```

module OHLCVTypes
  def self.included(base)
    base.coerce :date, Date
    base.coerce :close, Float
    base.coerce :volume, Integer
  end
end

```

```

    end
end

```

Dimension alignment coercion — map dimensions for * composition:

```

module TemporalAlignment
  def self.included(base)
    base.coerce :date, to: :quarter do |date|
      "#{date.year}-Q#{((date.month - 1) / 3) + 1}"
    end
  end
end
end

```

Coercion modules can be placed on the analysis class, on a loader, or on an ad hoc instance. Namo doesn't have an opinion about where they live. Inspired by Hashie's coercion extensions, but applied at the dimensional level.

The two forms are distinguishable by signature:

- `coerce :date, Date` — ingestion, fix the type of this dimension's values.
- `coerce :date, to: :quarter do ... end` — alignment, map this dimension to another for *.

Nested data flattening

Data from JSON APIs and other sources may arrive with nested structures:

```
{symbol: 'BHP', fundamentals: {eps: 1.2, pe: 14.5}}
```

Nested values break Namo's foundational principle — every key is a dimension. The solution is flattening at ingestion via a coercion declaration:

```
base.coerce :fundamentals, extract: [:eps, :pe, :book_value]
```

Which pulls specified keys up to the top level and drops the wrapper. The result is a flat hash where everything is a dimension.

Conversion discovery on *

When * finds no exact dimension name match, it checks the class's coercion registry for declared conversion paths. This replaces the earlier `respond_to?` design with explicit, inspectable, testable declarations.

Co-incident: Python bridge Paths 1 and 2 (separate repo)

Developed in a separate `namo-python` repository. The Python user writes Namo's Ruby DSL inside a triple-quoted string. The 3.x DSL block syntax is what they see:

```

from namo import Namo

analysis = Namo.define("""
  e.golden_cross   = proc{ sma(:close, 20) > sma(:close, 50) }
  e.earnings_yield = proc{ eps / close }
  e.action         = proc{ total_score >= 5 ? 'BUY' : total_score >= 2 ? 'WATCH' : 'HOLD' }
""")

results = analysis.run(data)

```

Path 1 (shell out): Python wrapper writes the DSL to a temp file, invokes ruby via subprocess, parses JSON results from stdout. Requires separate Ruby installation. ~0.5s overhead per invocation.

Path 2 (rb_call): MessagePack RPC to a persistent Ruby process. Lower latency (~50ms per call), more complex setup. Still requires Ruby installed separately.

Both paths build against 3.x Ruby Namo. The Python user experiences the full expressiveness of the DSL syntax. Ships with the Indicators module as a freebie so Python users get immediate value.

4.x: mRuby, SQL pushdown, and DuckDB

Theme: the database-integrated engine.

Estimated performance: DuckDB pushdown is the architectural win. Window functions (SMA, EMA) on millions of rows execute at ~0.3s in DuckDB versus ~45s in pure Ruby. The 2,000 row daily screen drops to ~0.05-0.1s total — DuckDB handles the heavy computation, Ruby handles the scoring and selection on the small result set. For the full 8.8M row ASX history, DuckDB makes it feasible where pure Ruby never could. Competitive with Polars (~0.1s) on equivalent workloads. All estimates.

mRuby compatibility

Namo's core compiled under mRuby. This enables embedding Namu in other environments without a full CRuby installation. Connects to existing mRuby compatibility work (e.g. `stateful.rb` 2.1.0).

DuckDB integration

DuckDB as a high-performance backend via Sequel or the `duckdb` gem. Four paths were considered, ordered from most manual to most automatic:

Path A (rejected): Write raw SQL for large slabs of functionality. Rejected — means maintaining SQL alongside Ruby, and the two representations can drift apart.

Path D (immediate): Use Sequel to generate DuckDB window functions before Namu ingests the data. No changes to Namu.

Path C (next): Formulae carry SQL metadata alongside the Ruby implementation. Namu detects the backend and inlines SQL where available:

```
module Indicators
  def sma(row, namo, field, period)
    window = namo[symbol: symbol, date: ..date].last(period)
    window.sum(&field) / window.count.to_f
  end

  SQL = {
    sma: ->(field, period) {
      "AVG(#{field}) OVER (PARTITION BY symbol ORDER BY date ROWS BETWEEN #{period - 1} PRECEDING AND CURRENT ROW)"
    }
  }
end
```

Path B (eventually): Namu generates SQL from formula definitions automatically.

Different Namos in the same composition can use different backends. OHLCV from DuckDB with SQL pushdown, fundamentals from Sequel, macro data from CSV. They compose via `*` regardless of origin.

SQL pushdown from formula metadata

Formulae declare their SQL equivalents. Namu uses the SQL version when a database backend is detected and falls back to Ruby when it can't. The Ruby implementation is the truth. The SQL is the optimisation.

Both are declared in the same place.

Co-incident: Python bridge Path 3 (separate repo)

With mRuby compatibility in 4.x, the Python bridge can embed the mRuby interpreter directly. `pip install namo` works with no separate Ruby dependency. Built in the `namo-python` repo against the 4.x mRuby-compatible core.

Near-zero call overhead — in-process, no serialisation. Cross-platform binary distribution needed (wheels per platform).

5.x: Streaming, transpilation, and genetic algorithms

Theme: the analytical platform.

Estimated performance: streaming eliminates the memory constraint — 8.8M rows never loaded simultaneously. Processing cost is per-window rather than per-dataset. Transpilation to SQL or Python/Polars delegates execution entirely to native engines — Nammo becomes a specification tool with near-zero runtime cost. GA support benefits from DuckDB's speed on each candidate evaluation. Performance is bounded by the backend, not by Nammo. All estimates.

Streaming / progressive computation

For large datasets (8.8M+ rows), process data in windows rather than loading everything. Rolling buffer per symbol, incremental indicator computation, results retained, raw data discarded.

A streaming Nammo stays unfrozen — it accepts new rows, invalidates caches on mutation, and reflects current state. The analytical case is a Nammo that gets frozen after setup.

Streaming performance vs incremental maintenance. Pure-live plus access-scoped materialisation (0.29.0) makes computing over an unfrozen, growing Nammo $O(n)$ per access — enough that streaming is *usable* in the core with no special class, as the Mutability note holds for correctness. But $O(n)$ -per-access is not optimal for a high-frequency feed: recomputing a column every tick is $O(n)$ where folding one row into running formula state — $EMA' = f(EMA, tick)$, a window advancing by one — is $O(1)$ or $O(window)$. Incremental maintenance is a distinct mechanism from any cache: it does not remember results, it maintains them, and it is the genuinely-more-optimal streaming design. A cross-access cache *with invalidation* is not that optimisation — under active streaming every append invalidates, so it degenerates to no cache. Because incremental maintenance requires formulae to expose an update rule (see *Formula opacity* below), it is a candidate for a commercial, streaming-optimised tier that extends the open core rather than living in it: the OSS floor stays correct, live, and non-catastrophic (access-scoped removes the $O(n^2)$); the incremental engine is the paid ceiling. This narrows OSS 2.x — with streaming performance absorbed by the 1.x floor, 2.x keeps freeze-gated caching, the DSL, and coherence, and incremental streaming (with live composition) moves to the commercial tier.

Transpilation

Nammo as a specification language that emits executable code in other targets:

- SQL (DuckDB, PostgreSQL) for production deployment.
- Python/Polars for handoff to Python teams.

Explore interactively in Ruby. Deploy as transpiled output. The exploration tool and the production artifact are the same object viewed from different angles.

Genetic algorithm support

Parameterised formulae enable GA-based strategy optimisation:

```

namo do |params|
  value_score -> { pe < params[:pe_strong] ? 2 : pe < params[:pe_good] ? 1 : 0 }
  action      -> { total_score >= params[:buy_threshold] ? 'BUY' : 'HOLD' }
end

```

Genome is a hash of parameters. Fitness is backtest return. Mutation/crossover operate on hashes. DuckDB evaluates each candidate at native speed.

Speculative / unversioned

These options are not scheduled. They become relevant when need and adoption justify the investment.

Rust acceleration via Rutie + Thermite

Rust accelerates specific hot paths within Ruby-side Namu. Ruby remains the host. The native code is an optimisation layer underneath the Ruby API. Formula resolution, `method_missing`, and DSL stay in Ruby.

Rutie provides Ruby bindings for Rust via macros. Thermite handles compilation and gem distribution with pre-compiled binaries. ~20x speedup observed on parsing benchmarks.

Targets: selection, iteration, storage layout. The same hot paths as the C extension (3.x) but with Rust's memory safety and the Rust ecosystem's momentum in data tooling (Polars, DuckDB, Apache Arrow).

```

namo/
  lib/
    namo.rb          # pure Ruby, always works
    namo/ext.rb      # loads native extension, falls back to pure Ruby
  ext/
    namo_ext/
      Cargo.toml     # Rust project
      src/
        lib.rs       # Rust implementations of hot paths
    Rakefile         # Thermite tasks for building

```

Work on Rust acceleration builds familiarity and infrastructure that would feed into the Rust-native engine if it ever happens.

FFI (generic)

Compile any language (Rust, C, Crystal) as a C-compatible shared library (`cdylib`), load via Ruby's `ffi` gem or `std::lib Fiddle`. More boilerplate than Rutie but language-agnostic — the same FFI interface works regardless of what produced the shared library.

Rust-native engine with Ruby DSL parser

Namu's core reimplemented in Rust with a parser for the Ruby DSL syntax. The DSL specification strings are parsed, not evaluated — no Ruby runtime at all. The Rust engine is language-agnostic and can be exposed to any host language via its native FFI mechanism:

- Python via PyO3 (the same way Polars works).
- Ruby via Rutie (replacing pure Ruby internals with Rust for speed).
- Node.js via napi-rs.
- Go via CGo.
- Any language via C-compatible FFI.

The Ruby DSL becomes a portable notation — a specification language that any host can submit to the Rust engine. The engine parses it, builds the internal representation, and executes it at native speed.

This path also connects to the transpilation story — if the Rust engine can parse the DSL, it can also emit SQL, Python/Polars expressions, or other output formats from the same parsed representation.

Not constrained to any single language. Largest engineering investment. Only justified if Namo achieves significant adoption across multiple language communities.

WASM compilation

Compile the Rust-native engine or mRuby core to WebAssembly for browser execution. Enables Namo-powered analysis in client-side web applications without a server.

Open questions

Mutability

Namo currently has `attr_accessor :data, :formulae`, making instances fully mutable. This conflicts with caching (stale results), thread safety (race conditions), and the principle that formulae are declarations.

The proposed model: Namos are mutable by default for interactive exploration. Call `freeze` when the shape is settled. Ruby's built-in `freeze` semantics apply — mutation methods (`[]=`, `.name=`) check `frozen?` and raise `FrozenError` if so. Caching is safe on frozen Namos. Threads can share frozen Namos.

```
namo = TradingAnalysis.new(ohlc_data)
namo.change = proc{ close - open }
namo.score = proc{ ... }

# exploration done
namo.freeze

# now cacheable, shareable, immutable
namo.score = proc{ ... } # => FrozenError
```

The streaming case is simply a Namo that never gets frozen. No special class needed — the distinction between analytical and streaming is a lifecycle difference, not a type difference. `freeze` is the transition point.

`attr_accessor` should become `attr_reader` on the public API, with mutation going through methods that check `frozen?`. Internal state (`@data`, `@formulae`) should be frozen when the Namo is frozen.

The `eql?/hash` contract introduced in 0.6.0 reinforces this. `eql?` requires hash to be consistent — `a.eql?(b)` implies `a.hash == b.hash` — and hash is computed from `@data`, `class`, and `@formulae`. Any mutation changes hash, which means an unfrozen Namo cannot be safely used as a Hash key or Set member: lookup might miss the entry the Namo was stored under because the Namo's hash has shifted. Frozen Namos have stable hashes and are safe for hash-keyed lookup. The lifecycle pattern (mutate during exploration, `freeze` before sharing) maps directly onto when `eql?`-based collection operations become reliable.

Live data and cache invalidation

For streaming/live Namos that stay unfrozen, formula results cached on Rows become stale when new data arrives. Options:

- No caching on unfrozen Namos (simple, slower).
- Generation counter — each mutation increments a counter, cached results carry the generation they were computed at, stale results are recomputed.
- Event-based invalidation — mutation notifies dependents.

The simplest approach (no caching when unfrozen) is probably right for the first implementation. Optimise if profiling shows it matters.

One structural fact bears on the choice: the deliberate `attr_accessor :data` back door (see the README's Mutability contract) means mutation can bypass every method Namo could instrument, so a generation counter cannot observe all mutation — which decides the 2.x mechanism in deep freeze's favour. Freezing `@data` when the Namo freezes closes the back door without removing the accessor, and caching then engages only where staleness is impossible by construction.

0.29.0 adds a fourth position the three options above missed. They all frame caching as *cross-access* — a result computed in one access, reused in a later one — and conclude “no caching when unfrozen.” Access-scoped materialisation caches only *within* one access and discards it on return: a parent-scanning formula asks for the same column once per row it visits, so materialising a column re-materialises its inputs n times over, an $O(n^2)$ that is pure redundancy inside a single access; coalescing it to one materialisation restores $O(n)$. This needs no invalidation at all. Nothing survives the access, so nothing can go stale, so there is nothing to observe — the back door is irrelevant, because a value that never outlives its access never needs a generation counter. It engages on unfrozen, streaming Namos, exactly where the freeze-gated 2.x cache cannot, and it is the same seam at a shorter lifetime: one access here, until-unfreeze there. The floor self-expires, the ceiling is immutable, and neither runs mutation-observing invalidation — so *no cache design in the roadmap needs to observe mutation*. The residue access-scoped does not touch — a formula genuinely rescanning a window per row — is not redundancy; it needs sorted/indexed windowed access, a separate primitive.

Formula opacity and the structured form

Four capabilities are each blocked by a formula being an opaque callable: by-value `===` (compares key sets instead — 0.6.0/0.23.0), dependency-aware live projection (materialise-and-drop instead — 0.16.0), portable serialisation (a proc's behaviour does not serialise), and incremental streaming maintenance (no update rule derivable from a black box). When one property blocks four capabilities, that is the signal it is a substrate concern, not a per-feature one. The bare-names DSL was scoped as projection's dependency enabler; the convergence argues for promoting a single *structured formula representation* to the privileged form the architecture organises around, with opaque procs kept as a supported escape hatch — runs in batch, forgoes these four.

The change to consider is a *widening*, not a *tightening*: structured becomes the primary form, opaque the fallback; opaque procs are never forbidden (some formulae are genuinely opaque — arbitrary Ruby, an external call — and cannot be reduced to a declarative form). This matches how Namo already treats formulae as a spectrum — arity dispatch is Namo inspecting a callable's shape to decide scope; “inspect its structure, when present, to unlock more” is the same move. The point is to design that representation once, rather than let projection, `===`, serialisation, and streaming each grow an incompatible notion of “structured.” Deciding the direction now costs nothing to build and stops new per-feature opacity workarounds accruing; the shipped ones (key-set `===`, materialise-and-drop) are not wasted — they become the behaviour of the opaque-escape tier. The 1.x releases and access-scoped materialisation (0.29.0) are contract-neutral — they neither need it nor foreclose it.

Coercion design

The coercion section (under 3.x) proposes a dual-purpose `coerce` keyword for both ingestion type fixing and dimensional alignment. Open questions:

- Is a single keyword for two different operations clear or confusing?
- Should ingestion coercion and alignment coercion be separate APIs?
- When multiple alignment coercions exist for a dimension (e.g. `:date` can coerce to both `:quarter` and `:month`), how does `*` choose? Raise ambiguity error? Accept a preference list?
- Should alignment coercion include a strategy (e.g. “most recent before”) or just a value transformation?
- How does coercion interact with `freeze`? Are coercion declarations frozen with the Namo?
- Should nested data flattening (`extract:`) be part of the coercion API or a separate ingestion concern?
- Should flattening be recursive (nested hashes within nested hashes) or single-level only?

Formula shadowing and method collisions

If a formula and a data dimension share the same name, which wins in Row resolution? Currently formulae take priority. Should this be configurable? Should it warn?

Related: Hashie::Mash's long history of problems with method names that collide with Ruby built-ins (`class`, `hash`, `sort`, `zip`, `count`) is directly relevant. Row's `method_missing` has the same risk — a dimension named `class` or `method` would shadow Ruby's built-in methods. Options:

- Maintain a safelist of Ruby method names that `method_missing` won't intercept.
- Warn when a dimension name collides with a built-in.
- Always allow hash-style access (`row[:class]`) as a fallback when bare name access is blocked.
- Document the risk and accept it — dimension names like `class` and `method` are unlikely in practice.

Operator return types with subclasses

When `TradingAnalysis * Namo` is evaluated, what class is the result? `TradingAnalysis` (preserving the included modules)? `Namo` (the base class)? The left operand's class? This affects whether formulae from included modules are available on the composed result.

0.6.0 settles part of this question for equality: `eq?` cares about class match (`TradingAnalysis.new(data).eq?(Namo.new(data))` returns false even if the data matches), `==` does not. 0.12.0's subclass guard pattern (`if name in initialize`, introduced with the `name: attribute`) addresses the side-effects-on-operator-results question — operator-derived instances are name-less and skip side effects. The class of operator results currently defaults to the receiver's class, which works for same-class composition. Cross-class composition (`TradingAnalysis * SectorMetrics`) still raises the question of which subclass's modules carry through.

Presentation examples

See `EXAMPLES.md` for full four-stage progressions (competitor tool → 1.x → 2.x → 3.x) across seven disciplines with side-by-side code comparisons.